

Jin-Soo Kim
(jinsoo.kim@snu.ac.kr)

Systems Software &
Architecture Lab.
Seoul National University

Fall 2019

Integers



Computer Systems



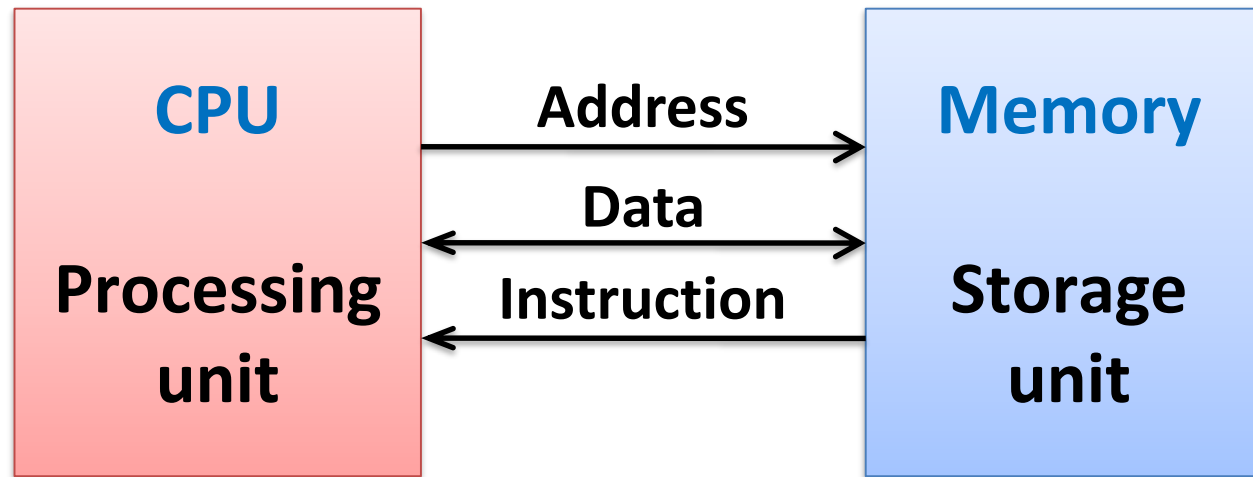
Google TV™



A computer is a **programmable** machine.

Von Neumann Architecture

- By John von Neumann, 1945



Data movement
Arithmetic & logical ops
Control transfer

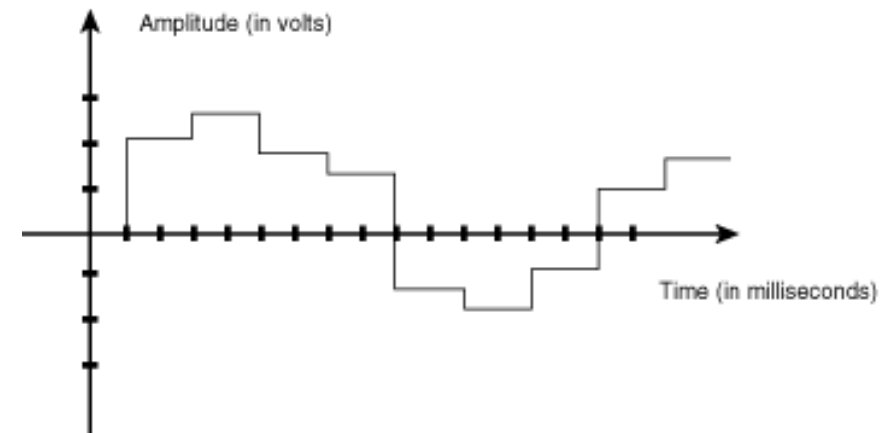

Byte addressable array
Code + data (user program, OS)
Stack to support procedures

The Advent of the Digital Age

- Analog vs. digital?

analog - continuous signal / storing the signal

<-> digital - digit(numbers) everything is converted into the numbers






- Compact Disc (CD)

- 44.1 KHz, 16-bit, 2-channel

- MP3 ^{frequency(주파수)}

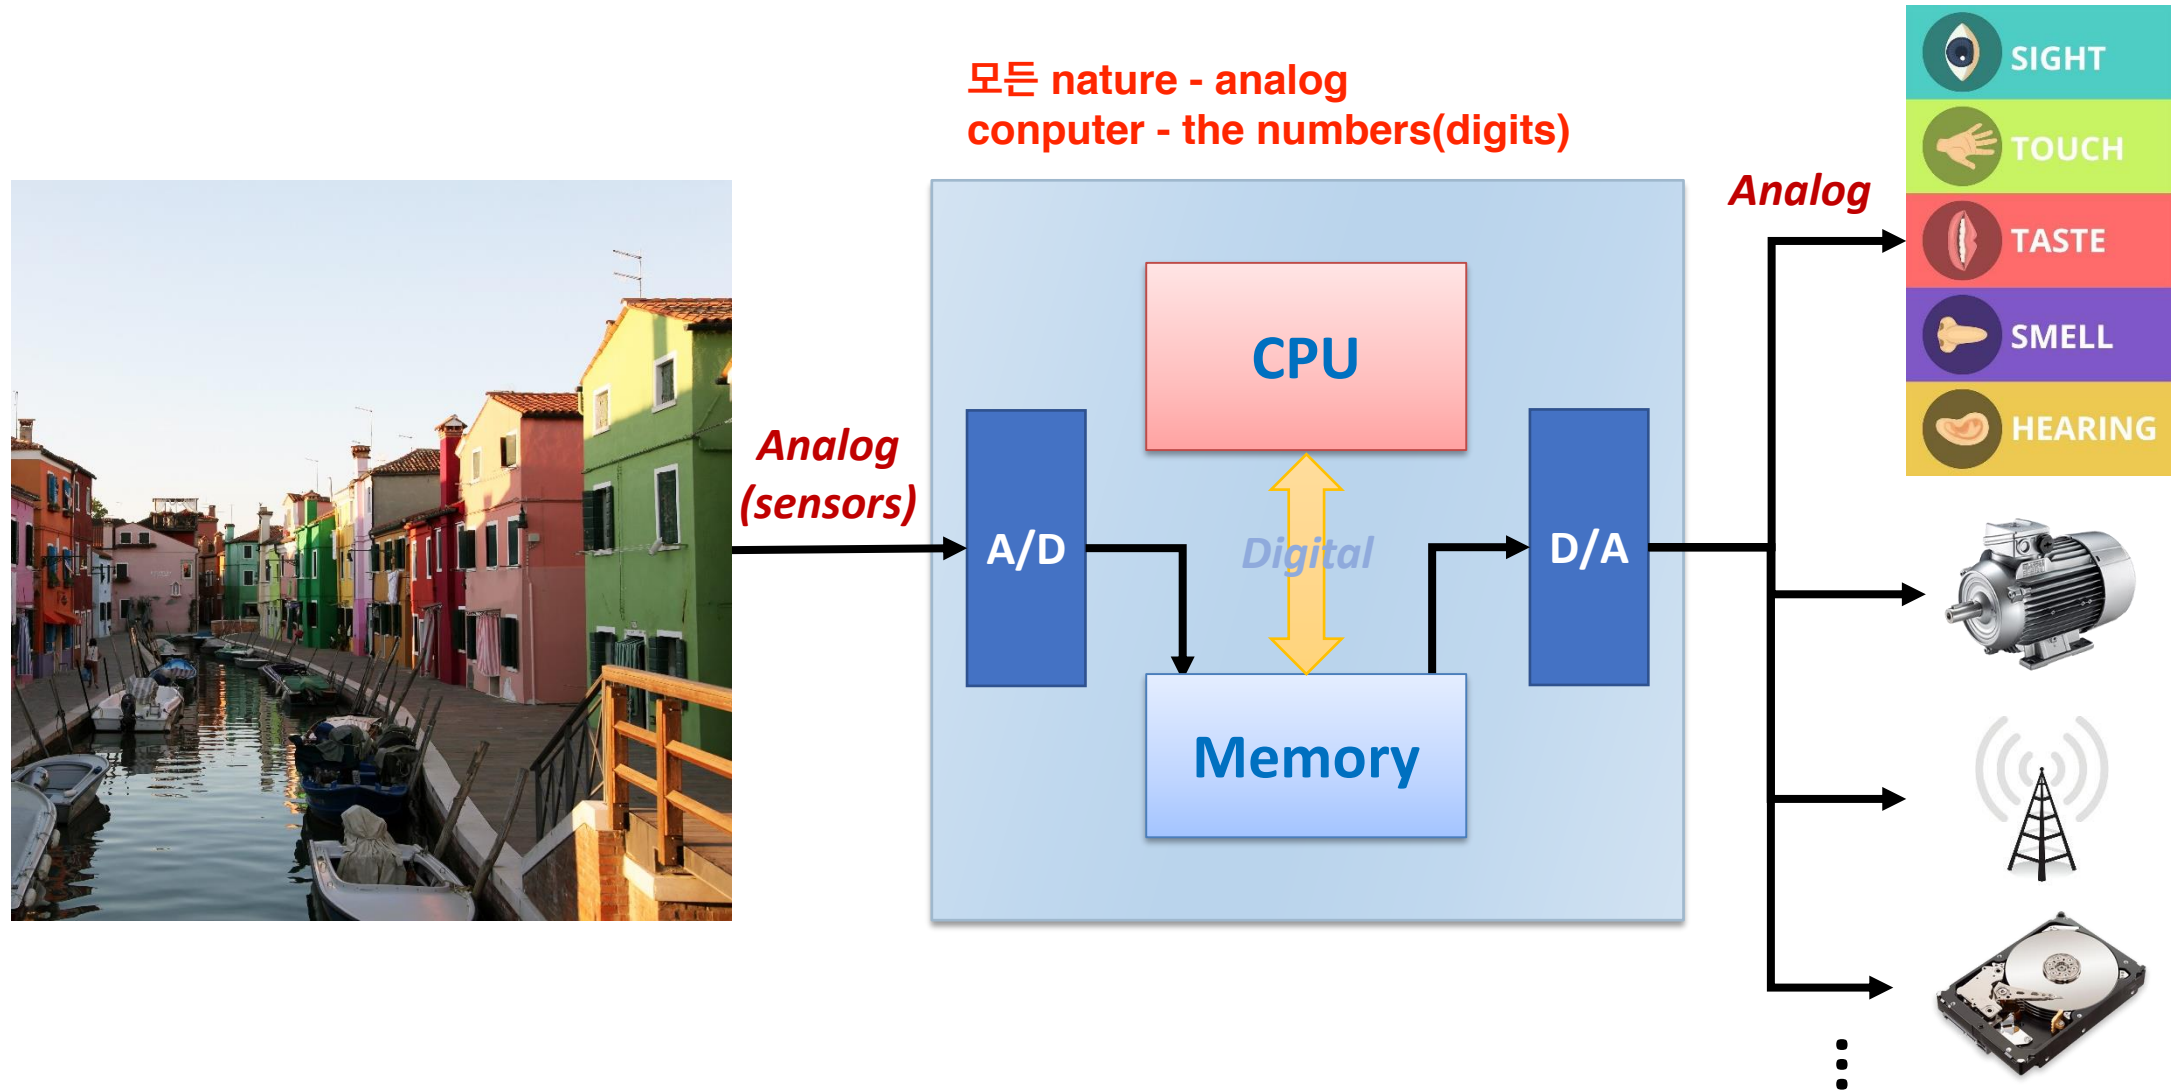
- A digital audio encoding with lossy data compression

LP Record vs. Compact Disc



Source: <http://www.soundsetal.com/blog-how-do-vinyl-records-work/>
<http://www.explainthatstuff.com/cdplayers.html>

Digital Computer



Representing Information

■ Information = Bits + Context


- Computers manipulate representations of things
- Things are represented as binary digits
- What can you represent with N bits?
 - 2^N things
 - Numbers, characters, pixels, positions, source code, executable files, machine instructions, ...
 - Depends on what operations you do on them

context 알아야 information 이해 가능

	01010011 01001110 01010101 01000001 01000010 01000011 01001000 01001001							
(char)	'S'	'N'	'U'	'A'	'R'	'C'	'H'	'I'
(int)	1096109651				1229472594			
(double)	1.08216479583800794... x 10 ⁴⁵							

Binary Representations

- Why not base 10 representation?
 - Easy to store with bistable elements **easy to store compute transfer**
 - Straightforward implementation of arithmetic functions
 - Reliably transmitted on noisy and inaccurate wires
- Electronic implementation



Encoding Byte Values

- Binary: 00000000_2 to 11111111_2
- Octal: 000_8 to 377_8
 - An integer constant that begins with 0 is an octal number in C
- Decimal: 0_{10} to 255_{10}
 - First digit must not be 0 in C
- Hexadecimal: 00_{16} to FF_{16}
 - Base 16 number representation
 - Use characters '0' to '9' and 'A' to 'F'
 - Write $FA1D37B_{16}$ in C as `0xFA1D37B` or `0xfa1d37b`

Hex	Decimal	Binary
0	0	0000
1	1	0001
2	2	0010
3	3	0011
4	4	0100
5	5	0101
6	6	0110
7	7	0111
8	8	1000
9	9	1001
A	10	1010
B	11	1011
C	12	1100
D	13	1101
E	14	1110
F	15	1111

Representing Integers

Unsigned Integers

- Encoding unsigned integers

$$B = [b_{w-1}, b_{w-2}, \dots, b_0] \quad x = 0001\ 0000\ 0101\ 1110_2$$

$$D(B) = \sum_{i=0}^{w-1} b_i \cdot 2^i$$

$$\begin{aligned} D(x) &= 2^{12} + 2^6 + 2^4 + 2^3 + 2^2 + 2^1 \\ &= 4096 + 64 + 16 + 8 + 4 + 2 \\ &= 4190 \end{aligned}$$


- What is the range for unsigned values with w bits? **2 to the $w - 1$ — maximum val**
- Using 64 bits: 0 to +18,446,774,073,709,551,615

Signed Integers

- Encoding positive numbers
 - Same as unsigned numbers
- Encoding negative numbers
 - Sign-magnitude representation
 - Ones' complement representation
 - Two's complement representation

Sign-magnitude Representation

- Two zeros
 - [000...00], [100..00]
- Used for floating-point numbers




$$S(B) = (-1)^{b_{w-1}} \cdot \left(\sum_{i=0}^{w-2} b_i \cdot 2^i \right)$$

2개의 0 0000 +0. 1000 -0.

$a-b > 0$ / $a-b == 0$

등등 여러 비교 시 +0/-0 사용에 문제




Ones' Complement Representation

- Easy to find $-n$
- Two zeros
 - $[000\dots00]$, $[111\dots11]$
- No longer used

just invert 0 / 1 flip!
2개의 0 문제점



$$O(B) = -b_{w-1}(2^{w-1} - 1) + \left(\sum_{i=0}^{w-2} b_i \cdot 2^i \right)$$




Two's Complement Representation (I)

- Unique zero
- Easy for hardware
 - leading 0 ≥ 0
 - leading 1 < 0
- Used by almost all modern machines



$$O(B) = -b_{w-1} \cdot 2^{w-1} + \left(\sum_{i=0}^{w-2} b_i \cdot 2^i \right)$$

1001(2)
-8. 1.
-8 + 1 = -7



Asymmetric range:

$$-2^{n-1} \sim 2^{n-1} - 1$$

Two's Complement Representation (2)

- Following holds for two's complement

$$\sim x + 1 == -x$$

- Complement

- Observation: $\sim x + x == 1111\dots11_2 == -1$

- Increment

$$\sim x + x == -1$$

$$\sim x + x + (-x + 1) == -1 + (-x + 1)$$

$$\sim x + 1 == -x$$

Sign Extension

short s : 16-bit
int i : 32-bit

s = i
i = s (filling with zero but, if it's minus


- Signed: w bits $\rightarrow w+k$ bits
 - Given w -bit signed integer x
 - Convert it to $w+k$ -bit integer with same value

- Replicate the sign bit to the left

- +2: 0000 0010 \rightarrow 0000 0000 0000 0010
- -2: 1111 1110 \rightarrow 1111 1111 1111 1110

unsigned는 상관없으니 0으로 채운다
type에 따라서 다르게 적용

- In RISC-V instruction set
 - lb: sign-extend loaded byte
 - lbu: zero-extend loaded byte



Manipulating Integers

Bit-Level Operations in C

- Operations `&`, `|`, `~`, `^` available in C
 - Apply to any “integral” data type
 - (unsigned) long, int, short, char
 - View arguments as bit vectors
 - Arguments applied bit-wise
- Examples (char data type)

$\sim 0x41 \rightarrow 0xBE$	$\sim 01000001_2 \rightarrow 10111110_2$
$\sim 0x00 \rightarrow 0xFF$	$\sim 00000000_2 \rightarrow 11111111_2$
$0x69 \ \& \ 0x55 \rightarrow 0x41$	$01101001_2 \ \& \ 01010101_2 \rightarrow 01000001_2$
$0x69 \ \ 0x55 \rightarrow 0x7D$	$01101001_2 \ \& \ 01010101_2 \rightarrow 01111101_2$
$0x69 \ ^ \ 0x55 \rightarrow 0x3C$	$01101001_2 \ \& \ 01010101_2 \rightarrow 00111100_2$

Logic Operations in C

- `&&`, `||`, `!`
 - View 0 as “False”, anything nonzero as “True”
 - Always return 0 or 1
 - Early termination
- Examples (char data type)

```
!0x41    →  0x00
```

```
!0x00    →  0x01
```

```
!!0x41   →  0x01
```

```
0x69 && 0x55 →  0x01
```

```
0x69 || 0x55 →  0x01
```

```
if (p && *p) ...           // avoids null pointer access
```

Shift Operations


- Left shift: $x \ll y$
 - Shift bit-vector x left y positions
 - Throw away extra bits on left
 - Fill with **0**'s on right
- Right shift: $x \gg y$
 - Shift bit-vector x right y positions
 - Throw away extra bits on right
 - Logical shift: fill with **0**'s on left
 - Arithmetic shift: replicate MSB on right
 - Useful with two's complement integer representation
- Undefined if $y < 0$ or $y \geq \text{word size}$

Argument x	01100010
$\ll 3$	00010000
Log. $\gg 2$	00011000
Arith. $\gg 2$	00011000

Argument x	10100010
$\ll 3$	00010000
Log. $\gg 2$	00101000
Arith. $\gg 2$	11101000

Addition

- Example: $7 + 6$



- Overflow if result out of range
 - Adding +ve and -ve operands: No overflow
 - Adding two +ve operands: Overflow if result sign is 1
 - Adding two -ve operands: Overflow if result sign is 0

Addition: Signed vs. Unsigned

- Signed addition in C
 - Ignores carry output
 - The low-order w bits are identical to unsigned addition

Examples for 3-bit integer

Mode	x	y	x + y	Truncated x + y
Unsigned	4 [100]	3 [011]	7 [0111]	7 [111]
Two's comp.	-4 [100]	3 [011]	-1 [1111]	-1 [111]
Unsigned	4 [100]	7 [111]	11 [1011]	3 [011]
Two's comp.	-4 [100]	-1 [111]	-5 [1011]	3 [011]
Unsigned	3 [011]	3 [011]	6 [0110]	6 [110]
Two's comp.	3 [011]	3 [011]	6 [0110]	-2 [110]

Subtraction


- Add negation of second operand
- Example: $7 - 6 = 7 + (-6)$

$$\begin{array}{r} +7: \quad 0000 \ 0000 \ \dots \ 0000 \ 0111 \\ -6: \quad 1111 \ 1111 \ \dots \ 1111 \ 1010 \\ \hline +1: \quad 0000 \ 0000 \ \dots \ 0000 \ 0001 \end{array}$$

- Overflow if result out of range
 - Subtracting two +ve or two -ve operands: No overflow
 - Subtracting +ve from -ve operand: Overflow if result sign is 0
 - Subtracting -ve from +ve operand: Overflow if result sign is 1

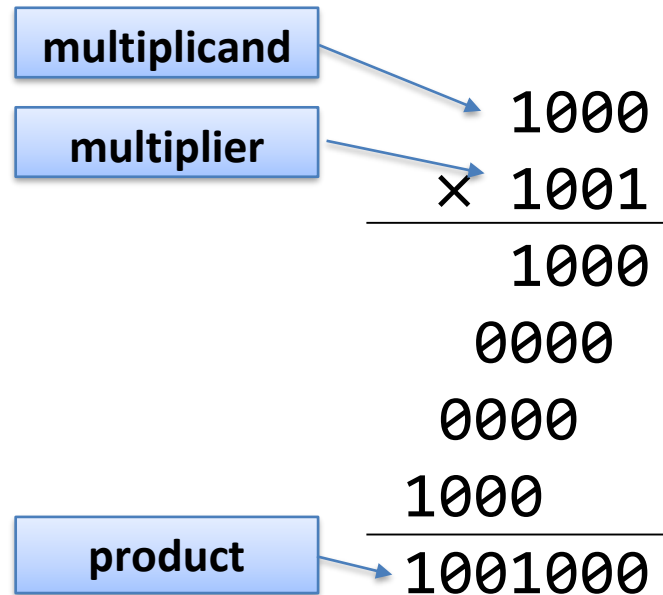
Arithmetic for Multimedia

- Graphics and media processing operates on vectors of 8-bit/16-bit data
 - Use 64-bit adder, with partitioned carry chain
 - Operate on 8x8-bit, 4x16-bit, or 2x32-bit vectors
 - SIMD (Single-Instruction, Multiple-Data)
- Saturating operations
 - On overflow, result is largest representable value
 - e.g., clipping in audio, saturation in video

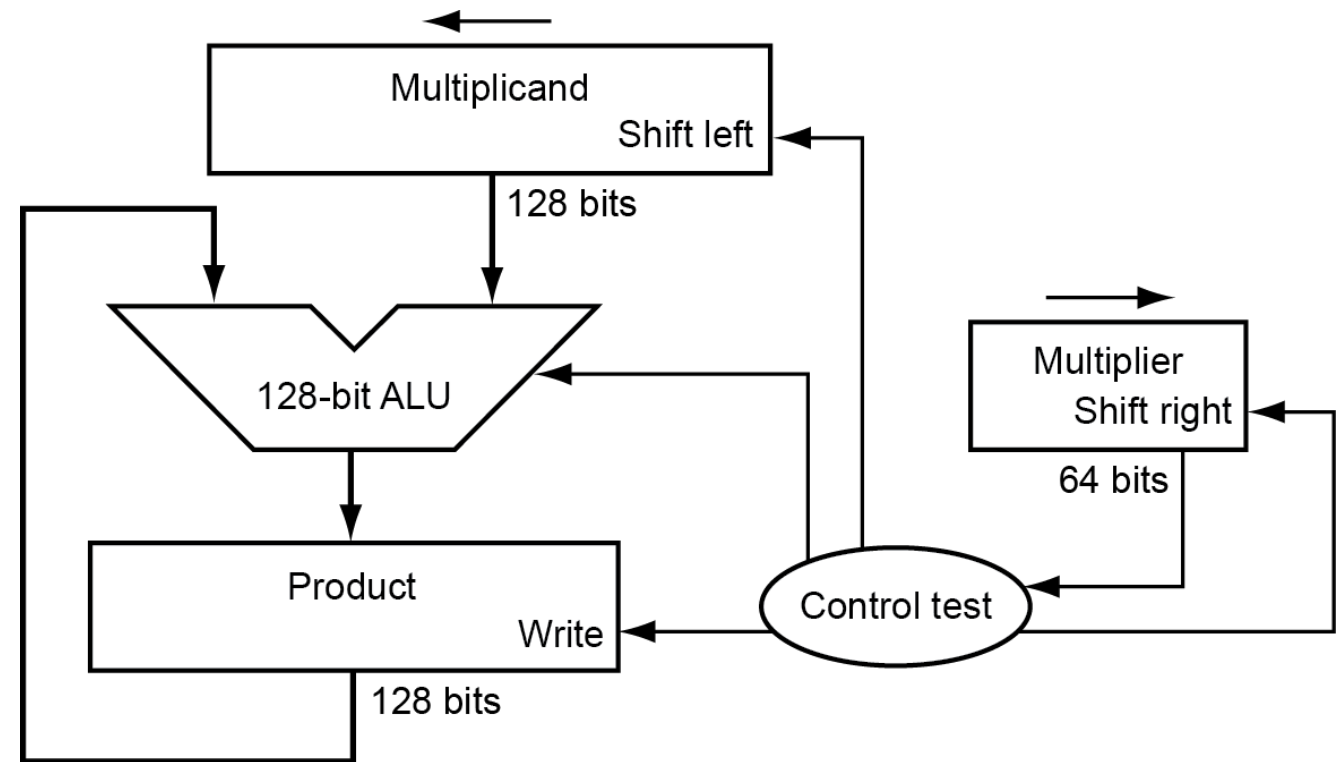


Multiplication

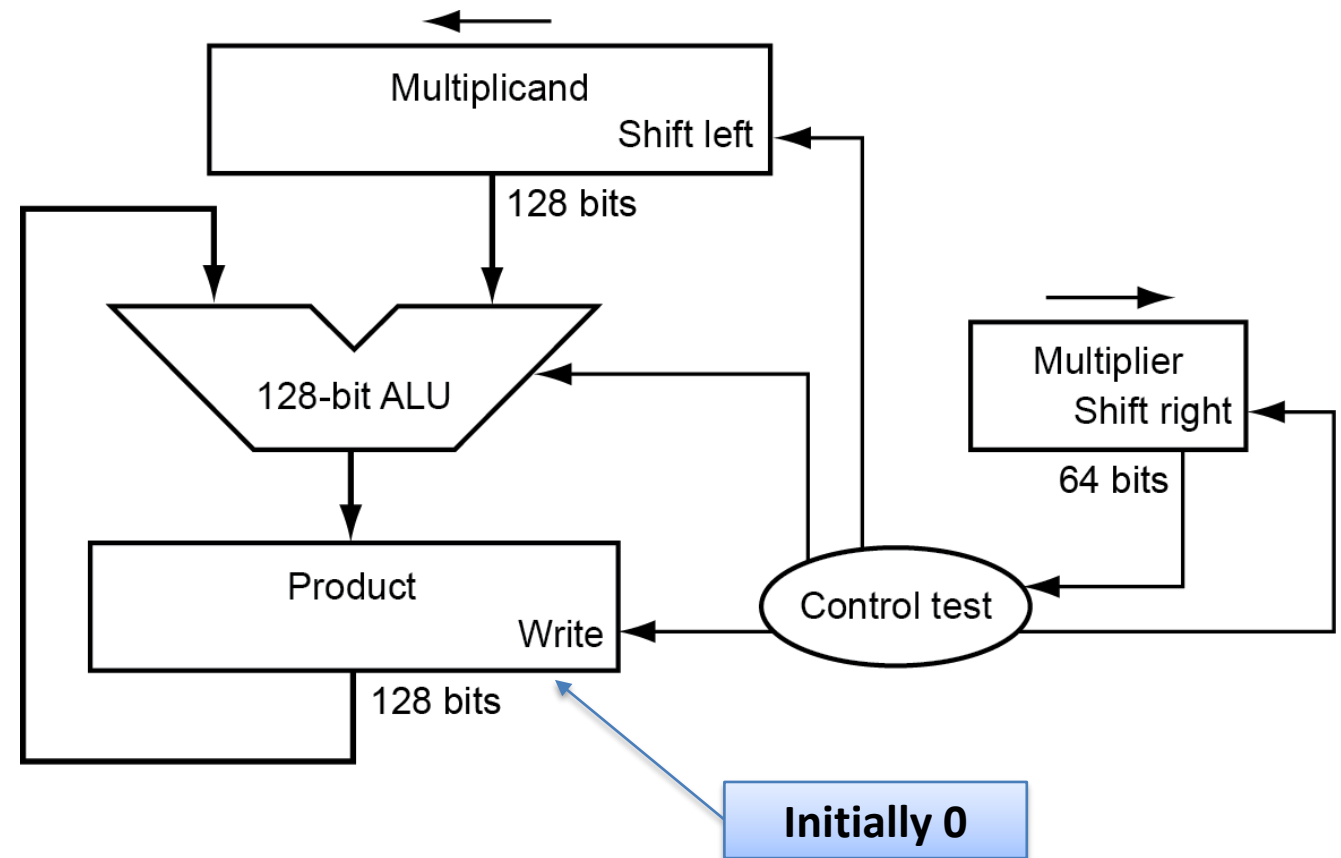
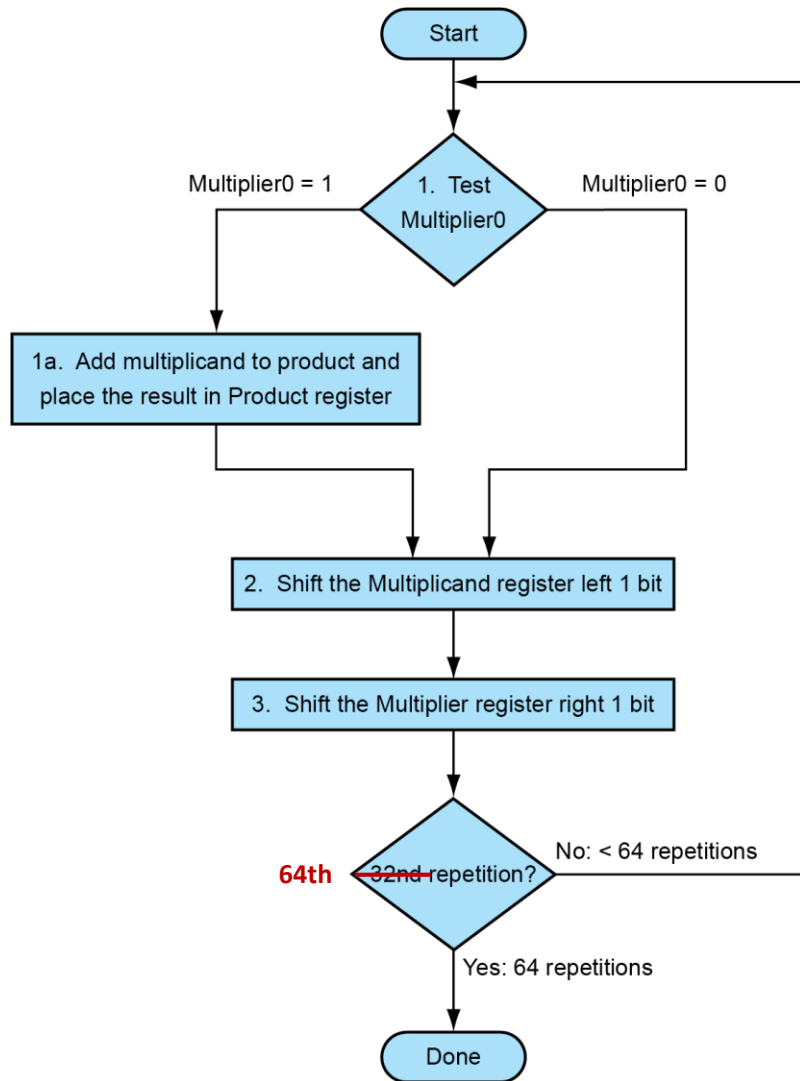
- Long-multiplication approach



Length of product is
the sum of operand
lengths

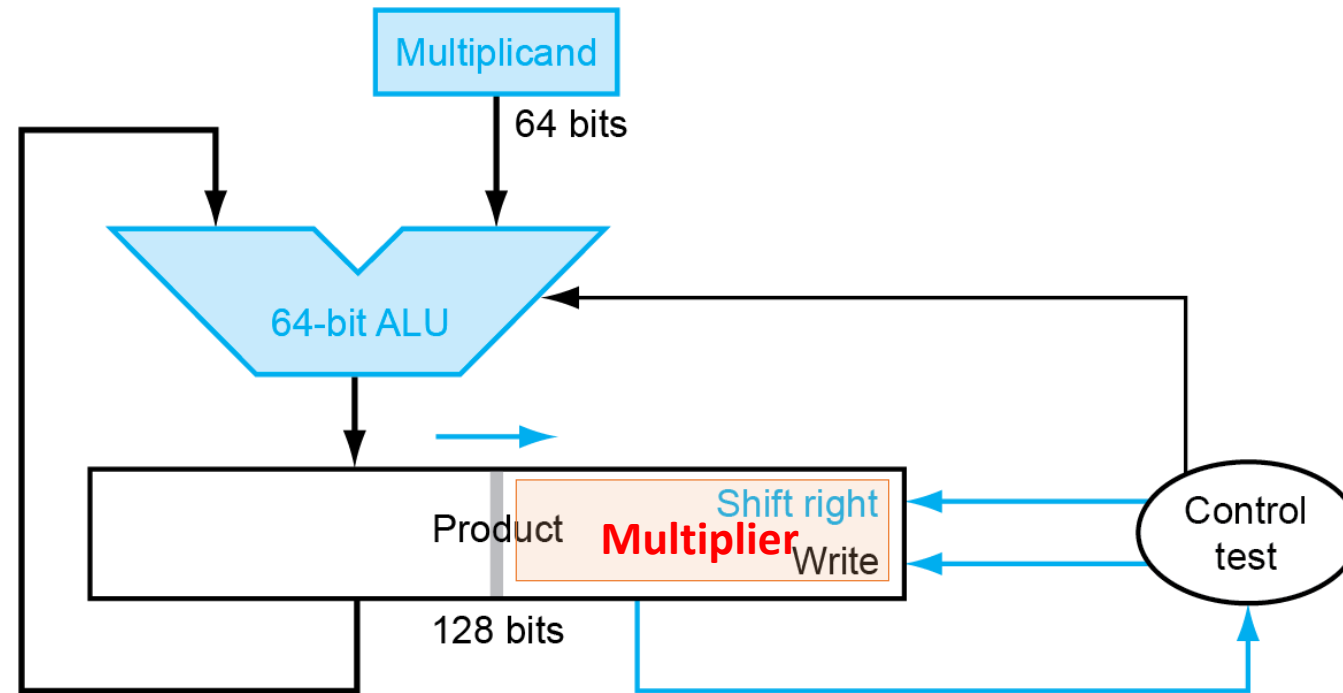


Multiplication Hardware



Optimized Multiplier

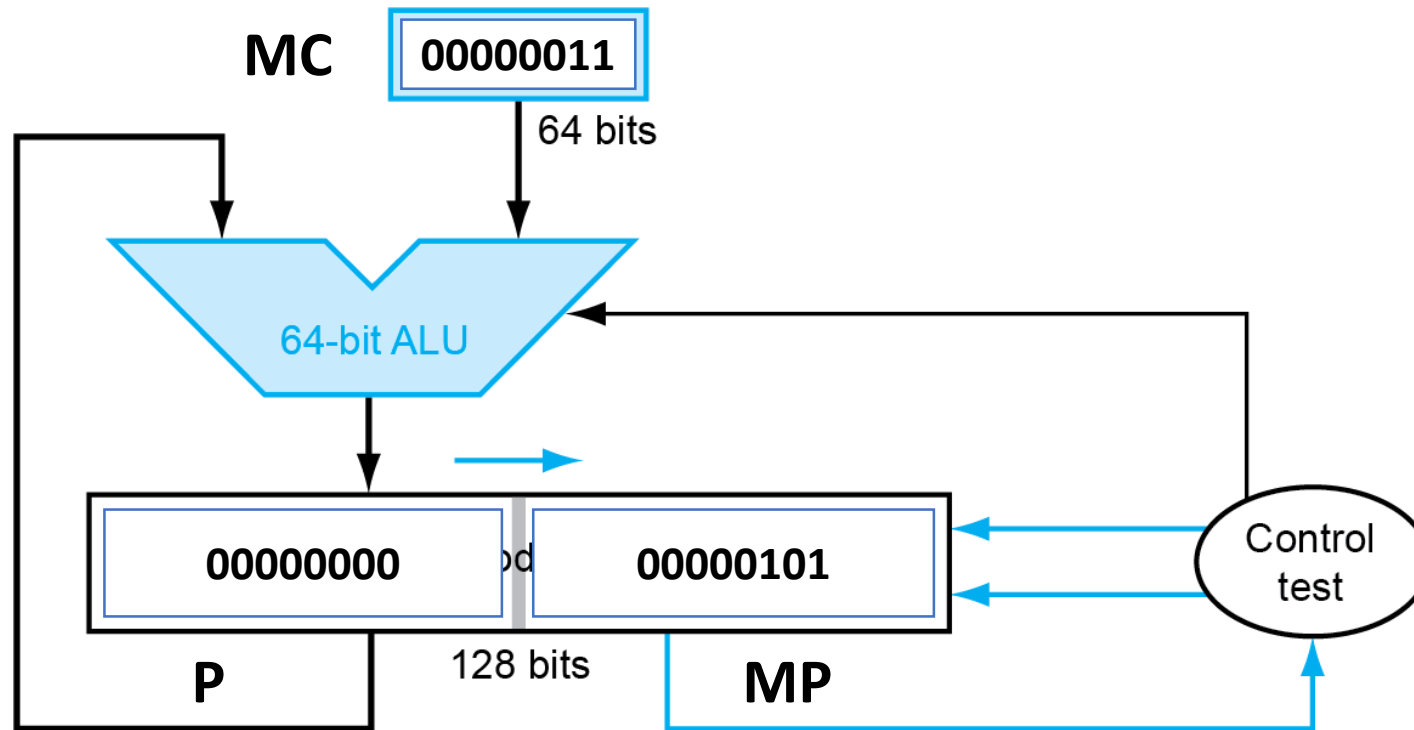
- Perform steps in parallel: add / shift



- One cycle per partial-product addition
 - That's ok, if frequency of multiplication is low

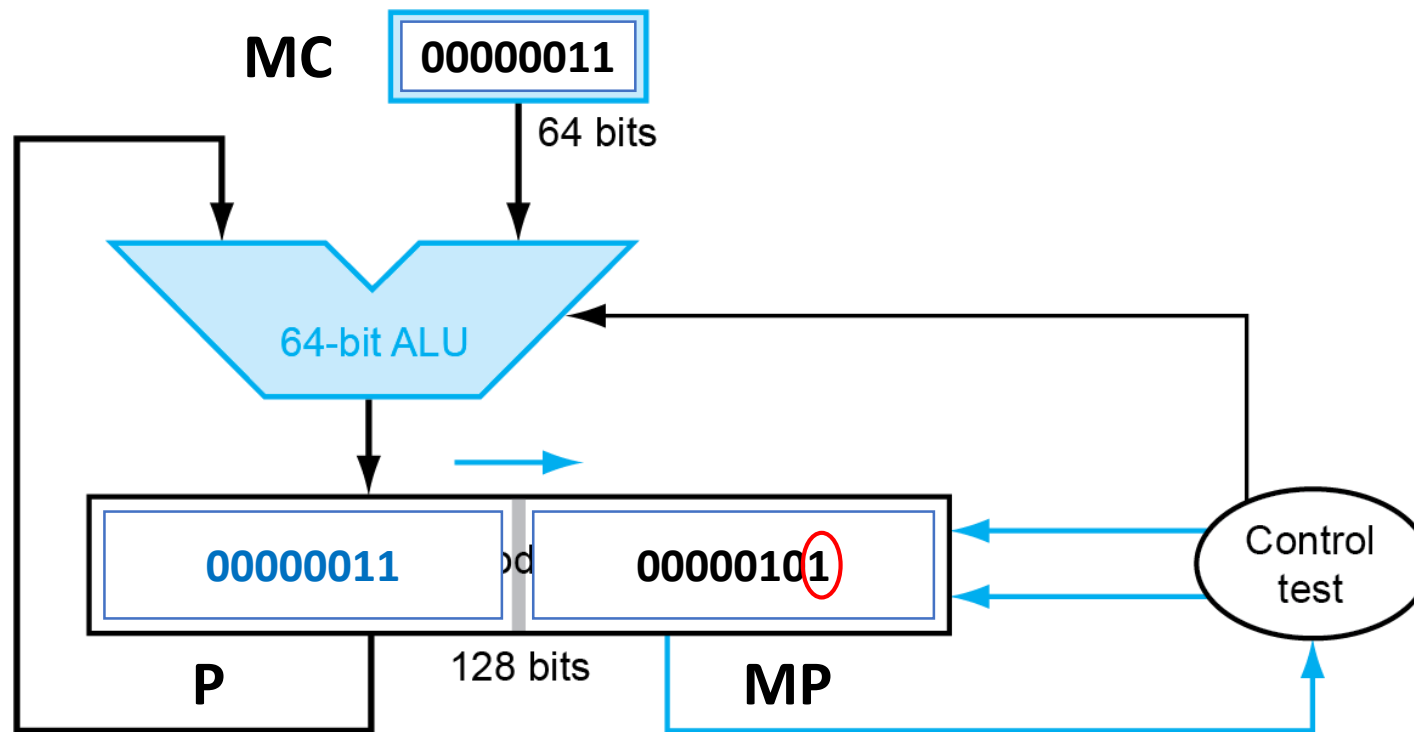
Multiplication Example (I)

▪ $3_{10} \times 5_{10} = 00000011_2 \times 00000101_2$



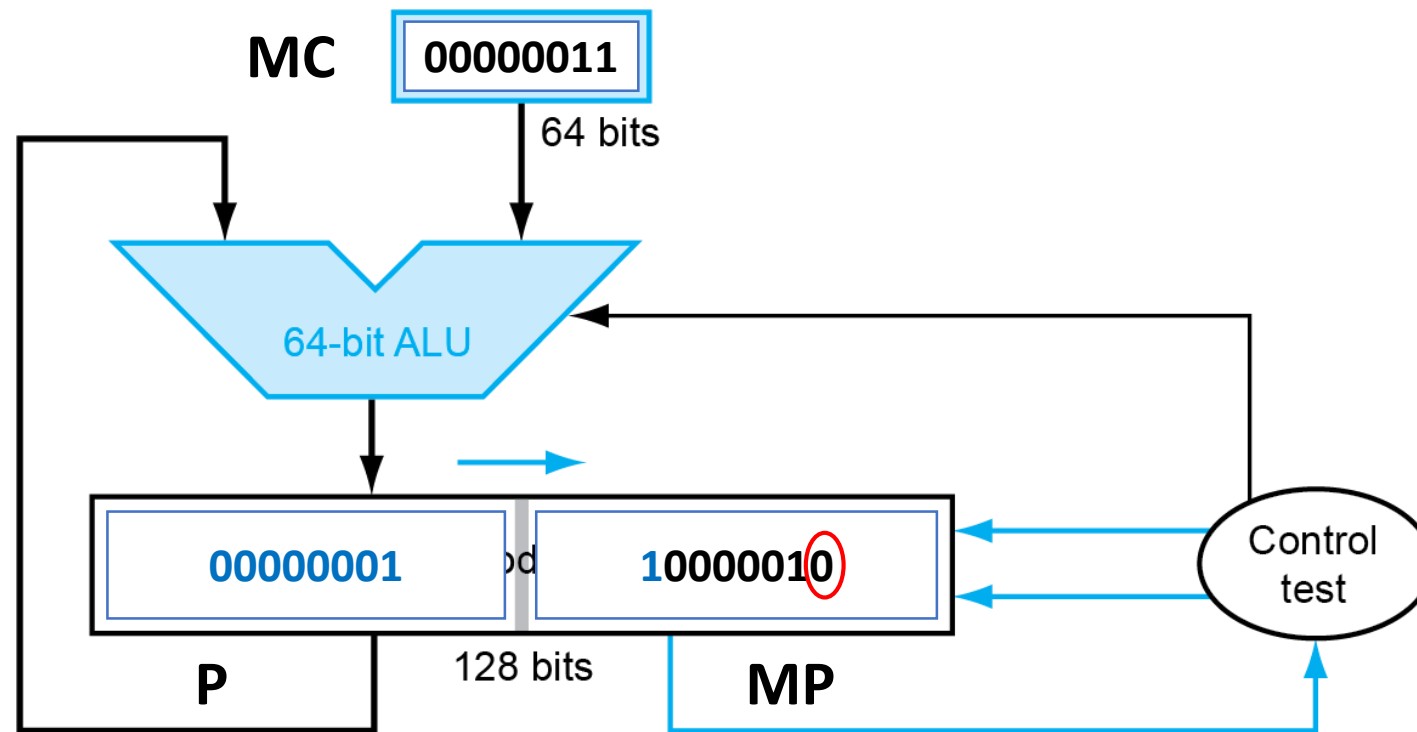
Multiplication Example (2)

- $MP_0 = 1$: Add MC to P, shift right P and MP by 1 bit



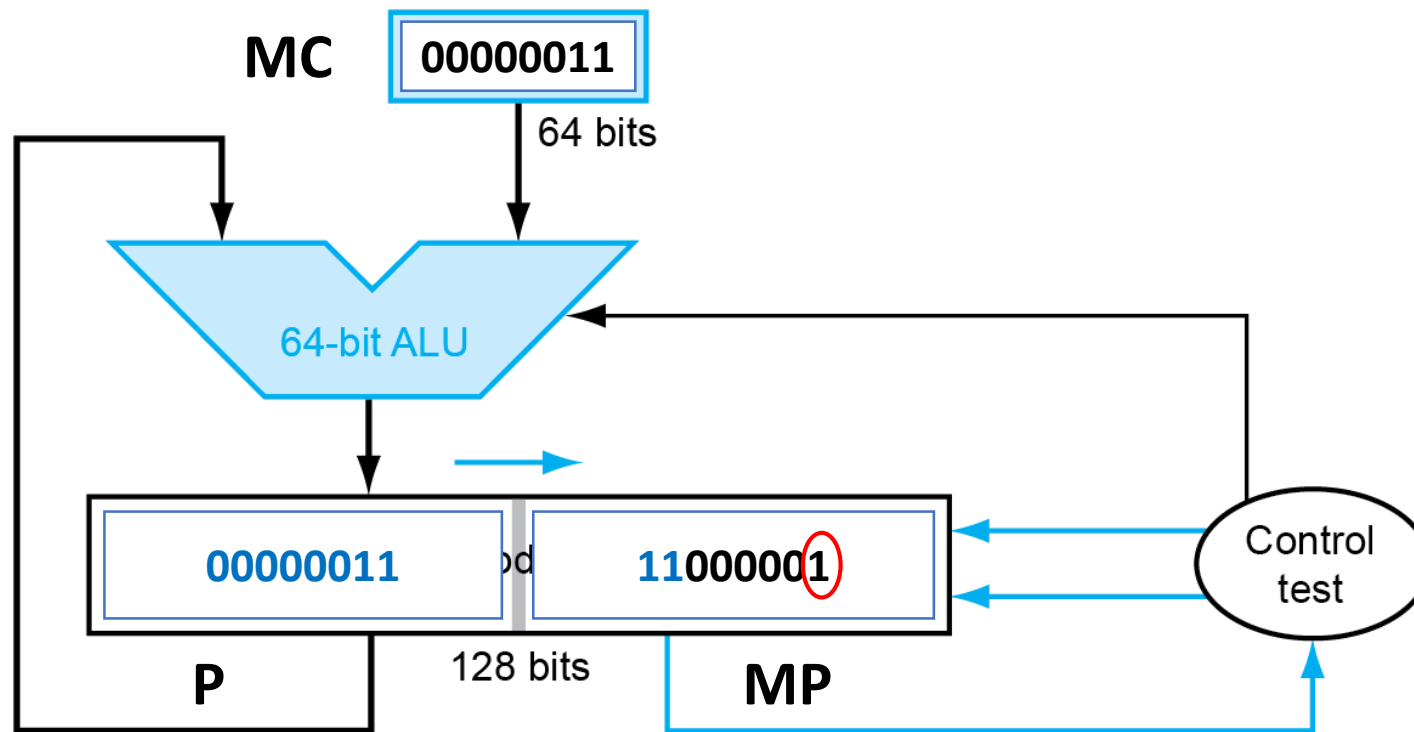
Multiplication Example (3)

- $MP_0 = 0$: Shift right P and MP by 1 bit



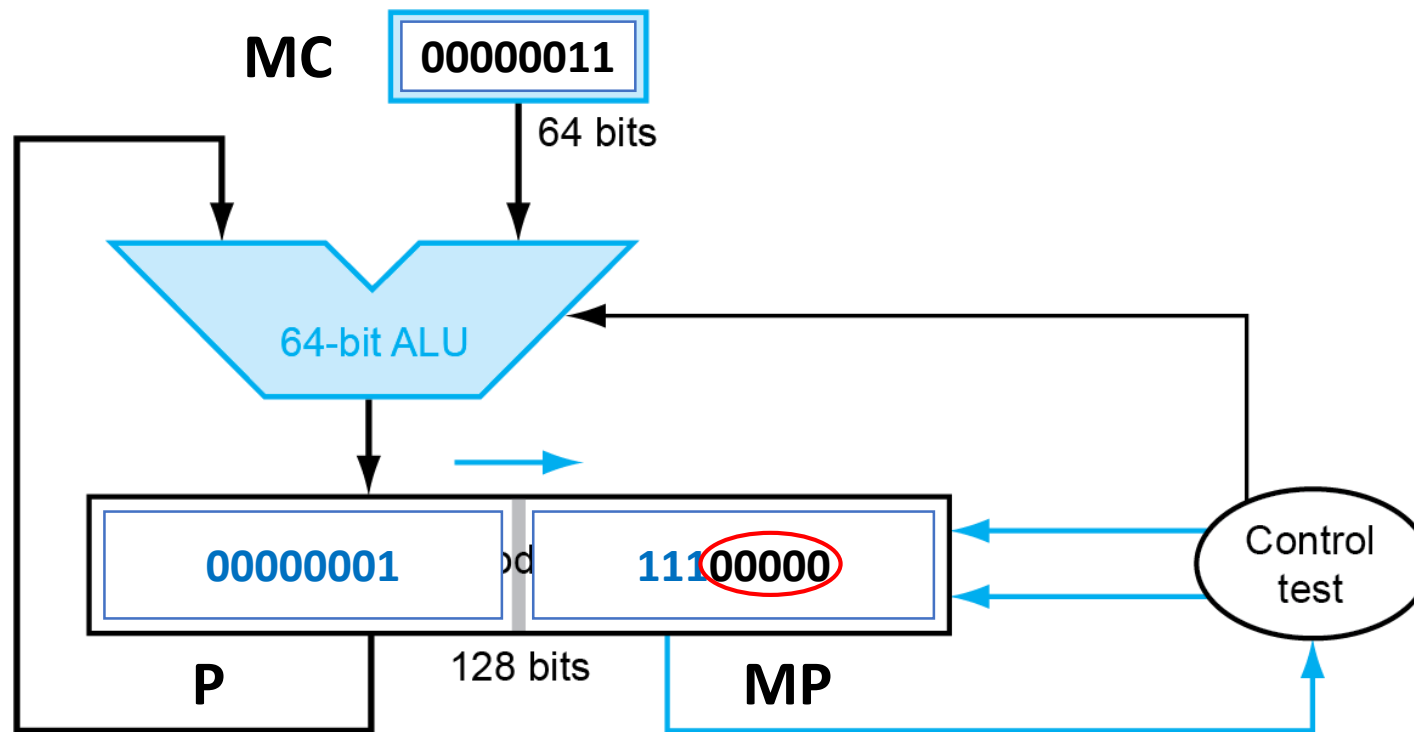
Multiplication Example (4)

- $MP_0 = 1$: Add MC to P, shift right P and MP by 1 bit



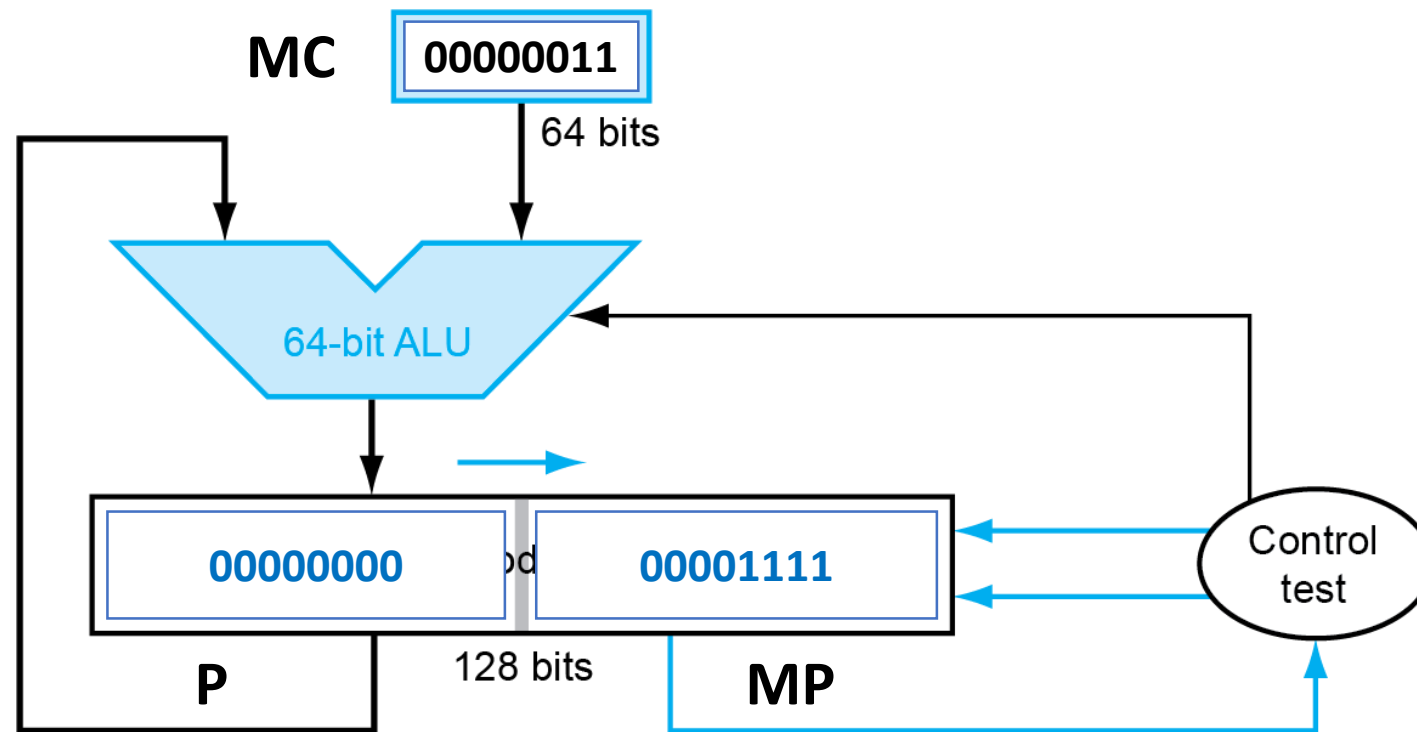
Multiplication Example (5)

- $MP_0 = 0$ (5 times): Shift right P and MP by 1 bit (5 times)



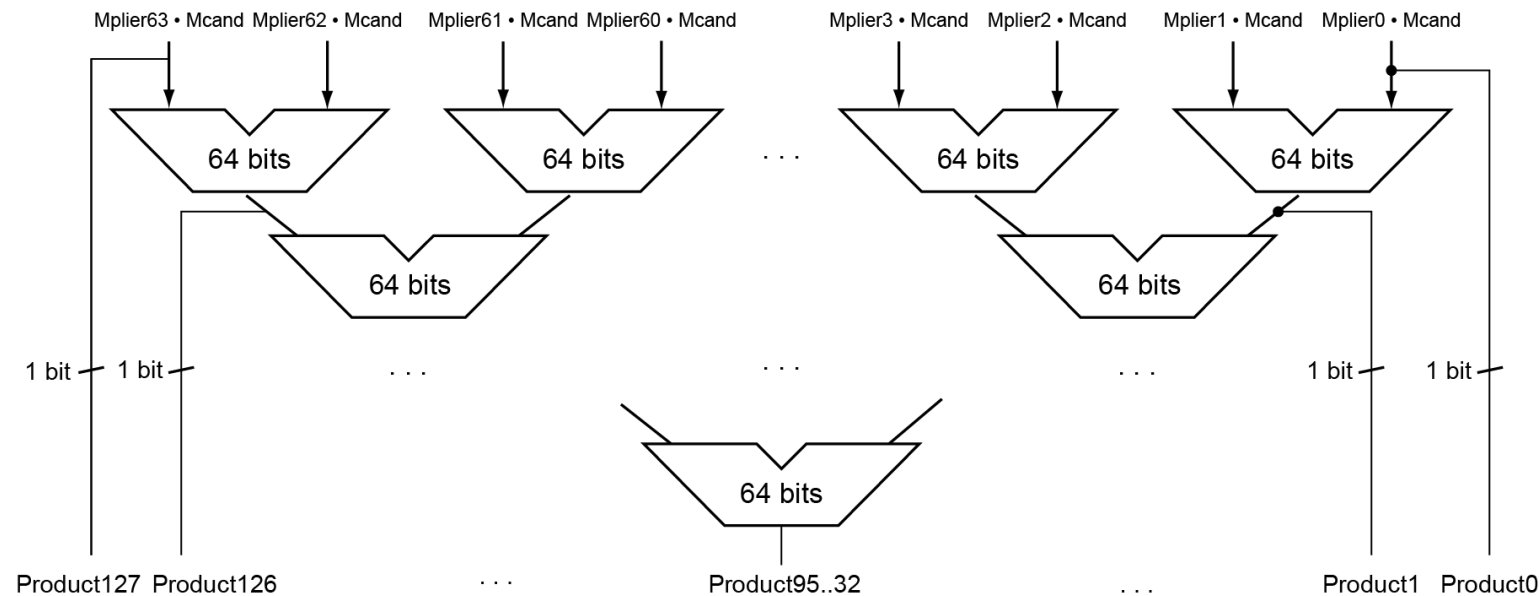
Multiplication Example (6)

- Final product available in P + MP



Faster Multiplier

- Uses multiple adders
 - Cost/performance tradeoff



- Can be pipelined
 - Several multiplication performed in parallel

Multiplication: Signed vs. Unsigned

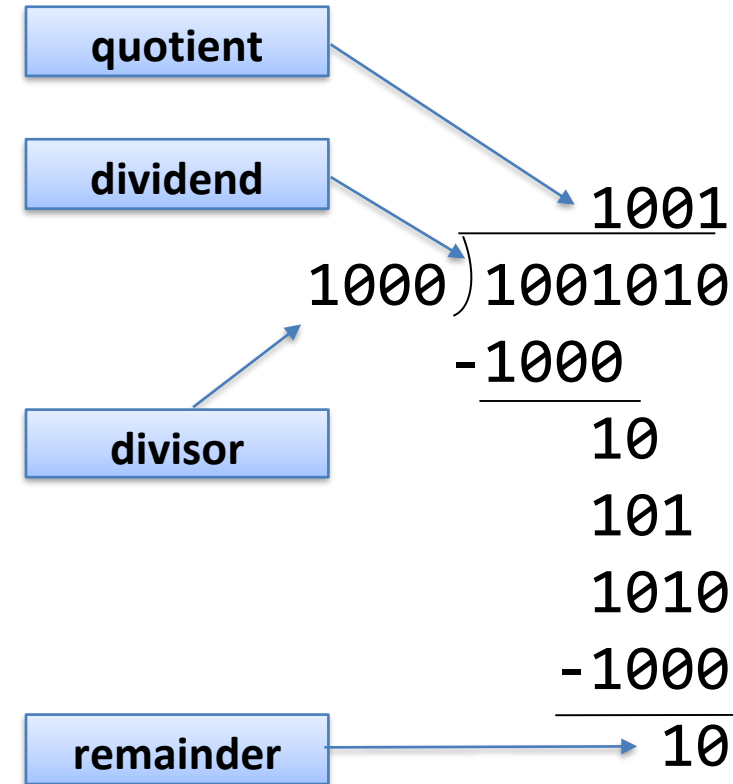
- Signed multiplication in C
 - Ignores high order w bits
 - The low-order w bits are identical to unsigned multiplication

Examples for 3-bit integer

Mode	x	y	$x \cdot y$	Truncated $x \cdot y$
Unsigned	5 [101]	3 [011]	15 [001111]	7 [111]
Two's comp.	-3 [101]	3 [011]	-9 [110111]	-1 [111]
Unsigned	4 [100]	7 [111]	28 [011100]	4 [100]
Two's comp.	-4 [100]	-1 [111]	4 [000100]	-4 [100]
Unsigned	3 [011]	3 [011]	9 [001001]	1 [001]
Two's comp.	3 [011]	3 [011]	9 [001001]	1 [001]

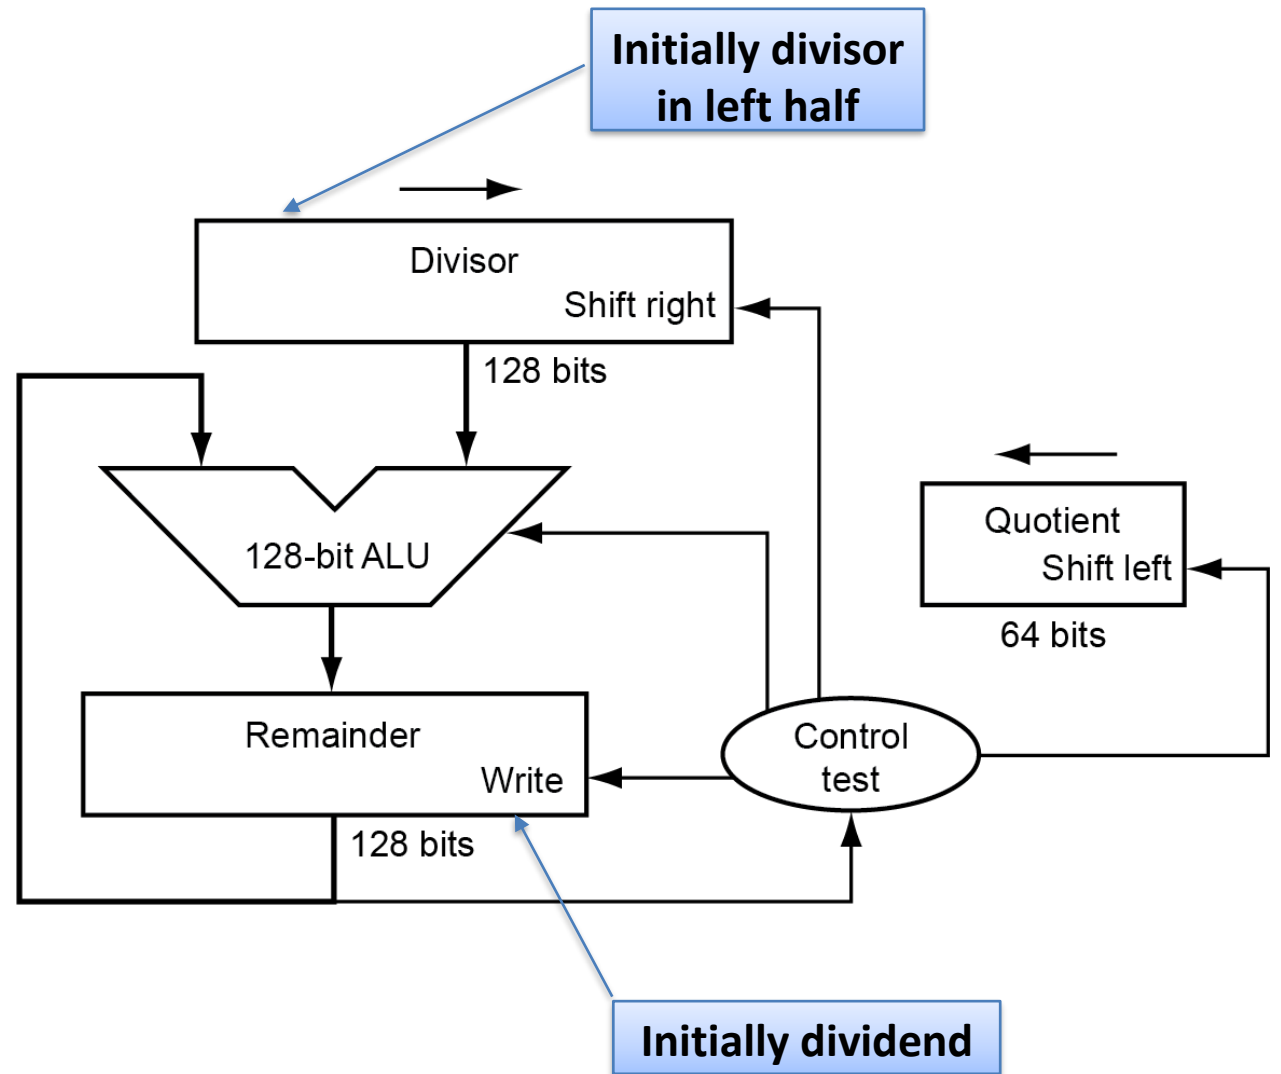
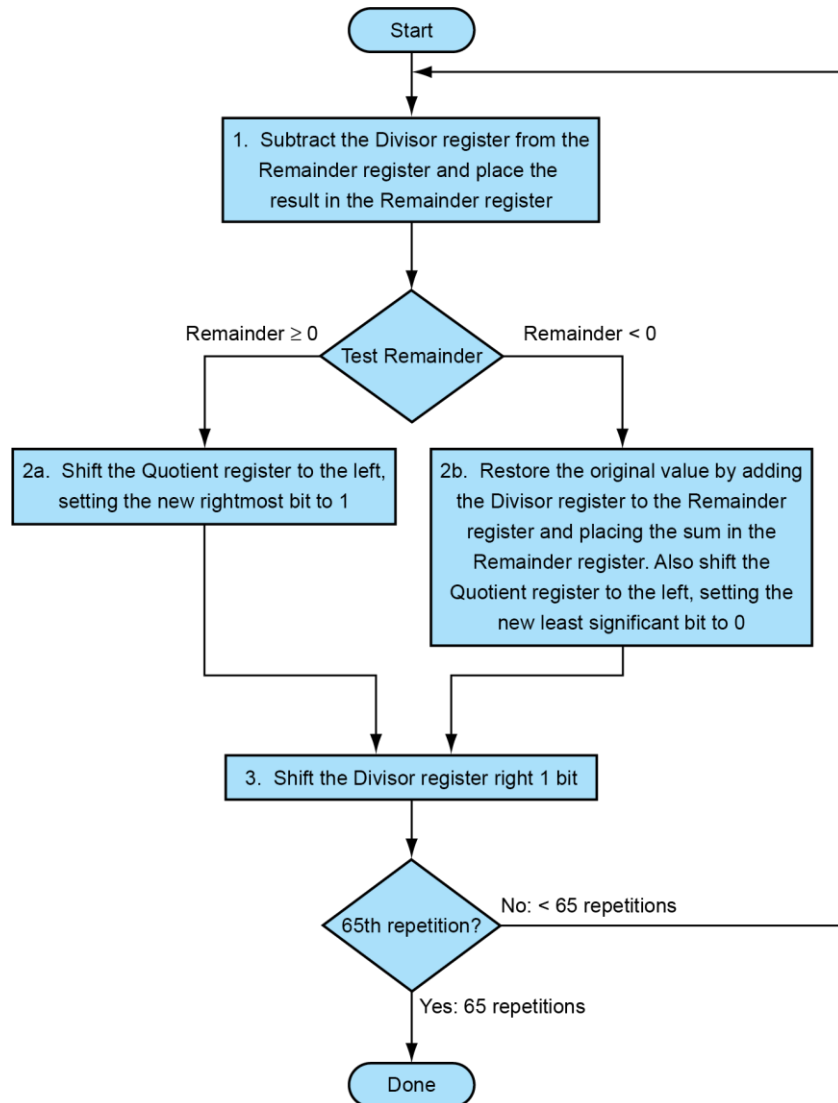
Division

- Check for 0 divisor
- If divisor \leq dividend bits:
1 bit in quotient, subtract
- Otherwise: 0 bit in quotient,
bring down next dividend bit
- Restoring division
 - Do the subtract, and if remainder goes < 0 , add divisor back

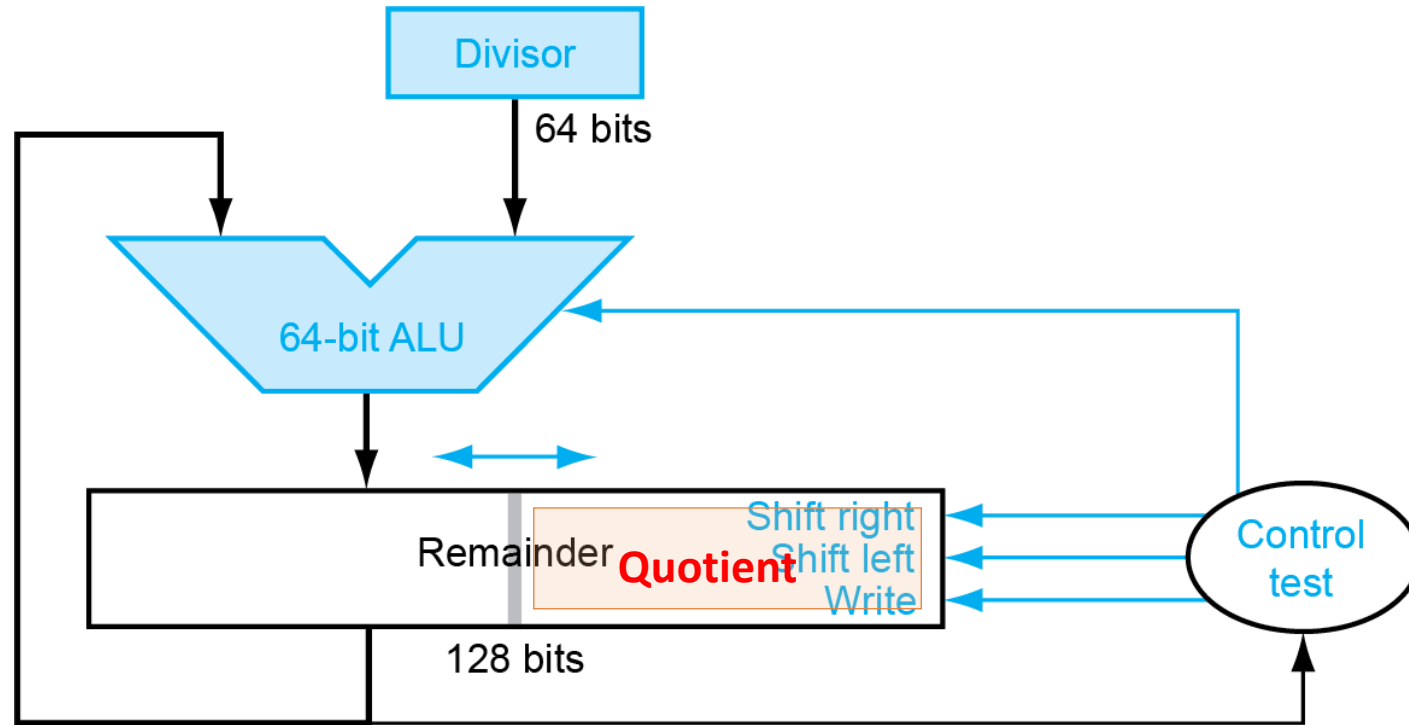


***n*-bit operands yield *n*-bit
quotient and remainder**

Division Hardware



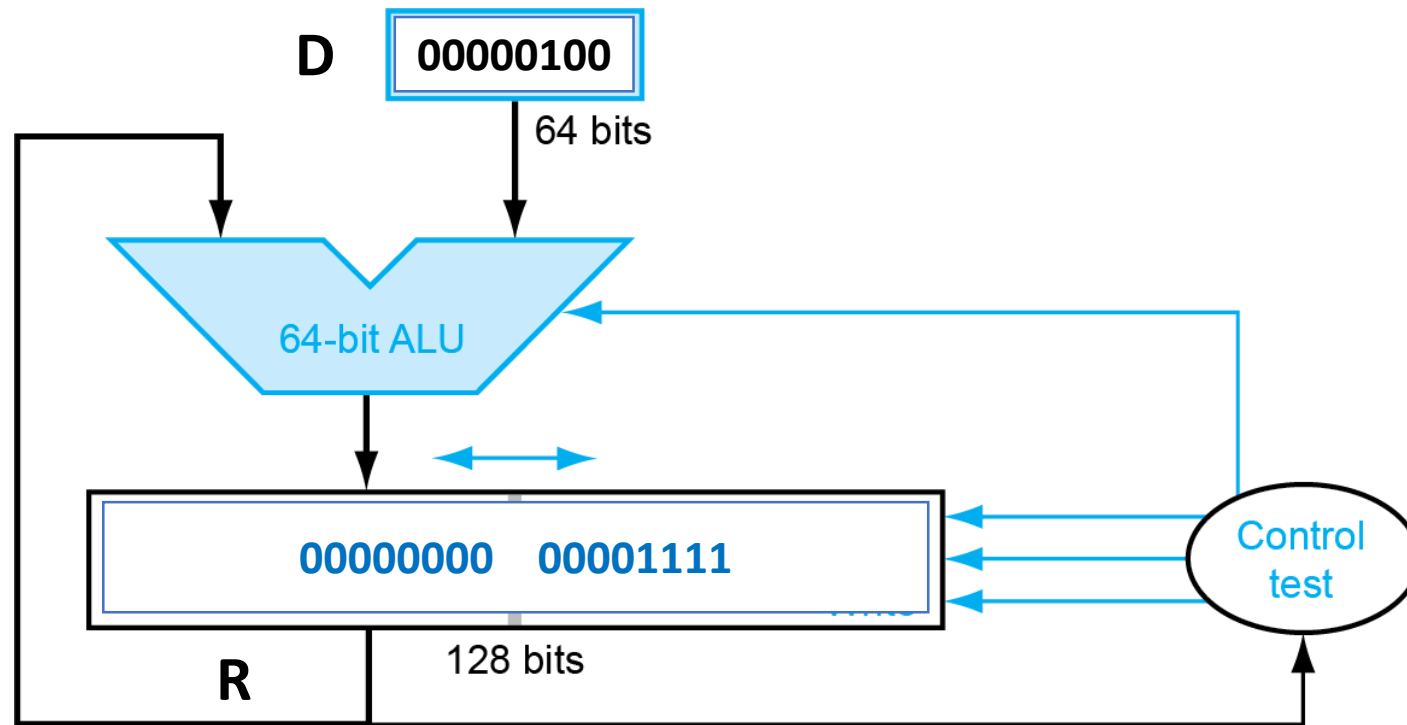
Optimized Divider



- One cycle per partial-remainder subtraction
- Looks a lot like a multiplier!
 - Same hardware can be used for both

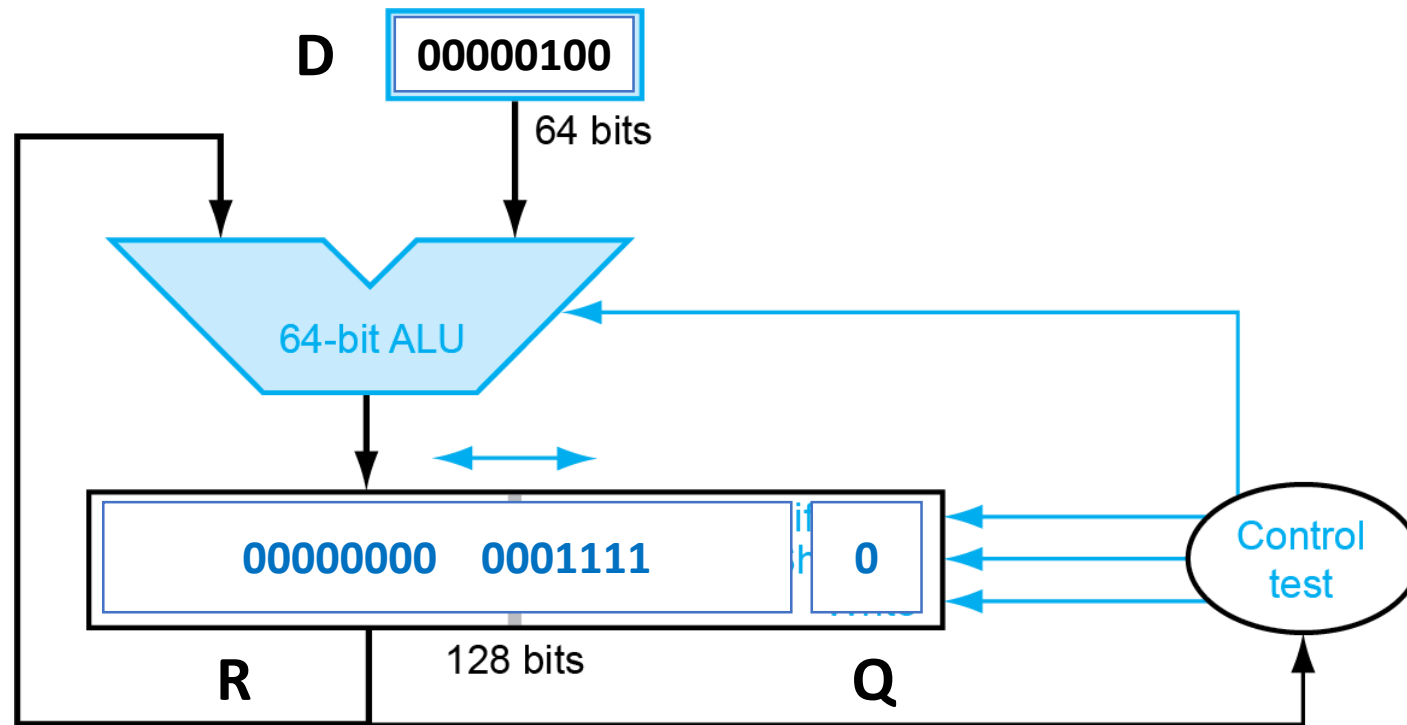
Division Example (I)

■ $15_{10} / 4_{10} = 00001111_2 / 00000100_2$



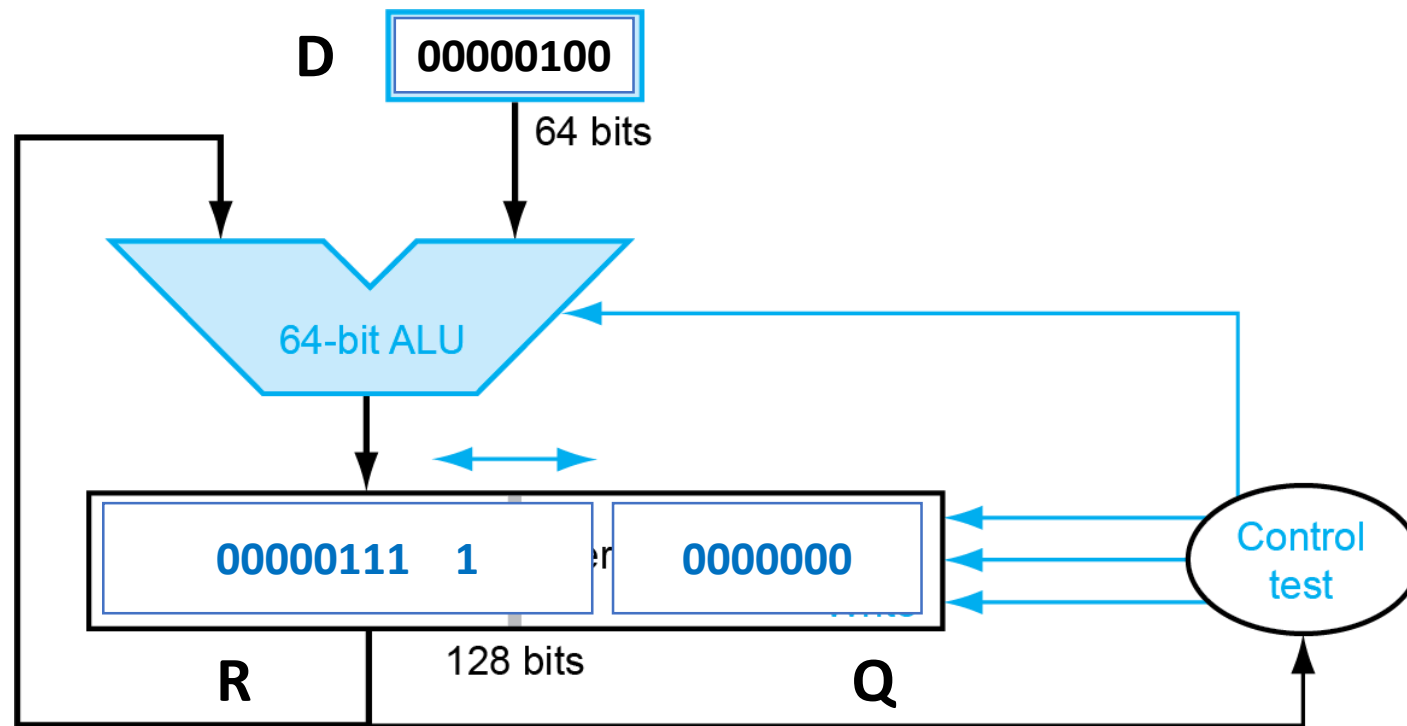
Division Example (2)

- $R - D < 0$: Shift left R and Q by 1 bit, $Q_0 = 0$



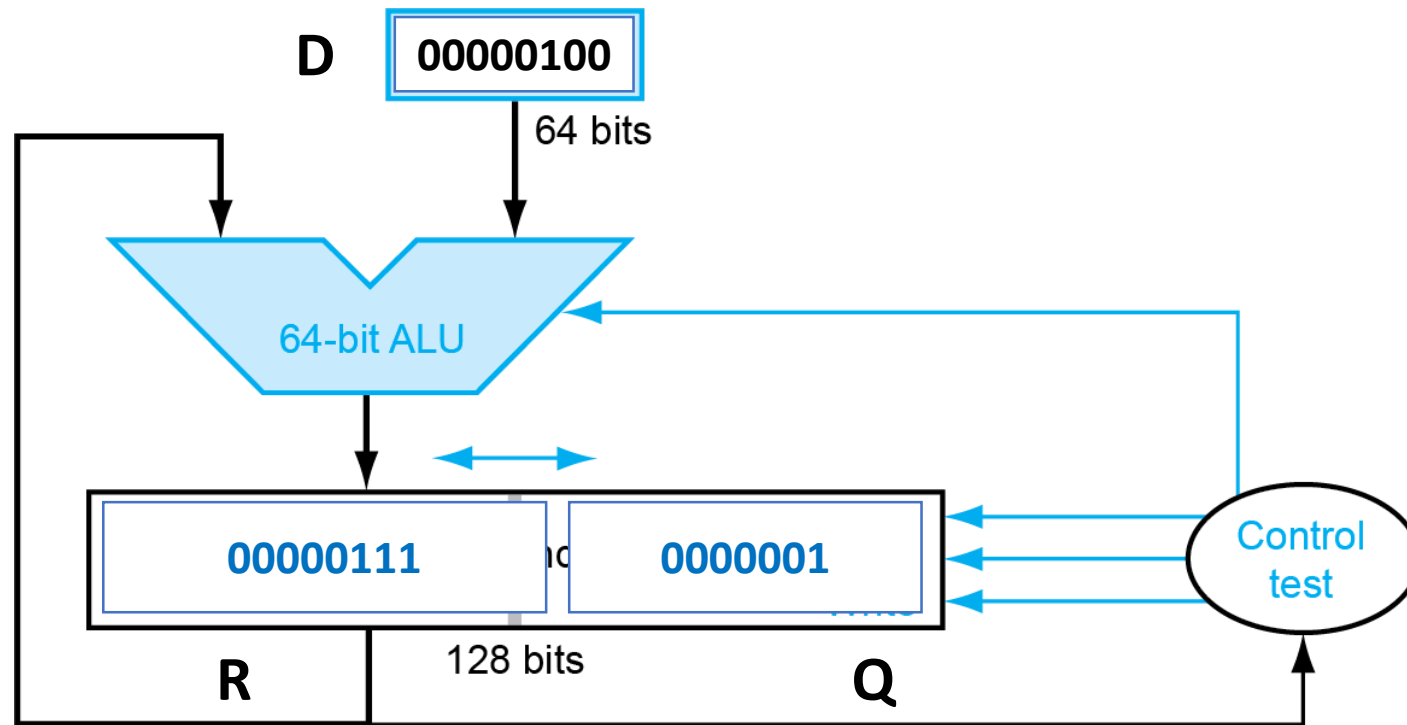
Division Example (3)

- $R - D < 0$: Shift left R and Q by 1 bit, $Q_0 = 0$ (6 times)



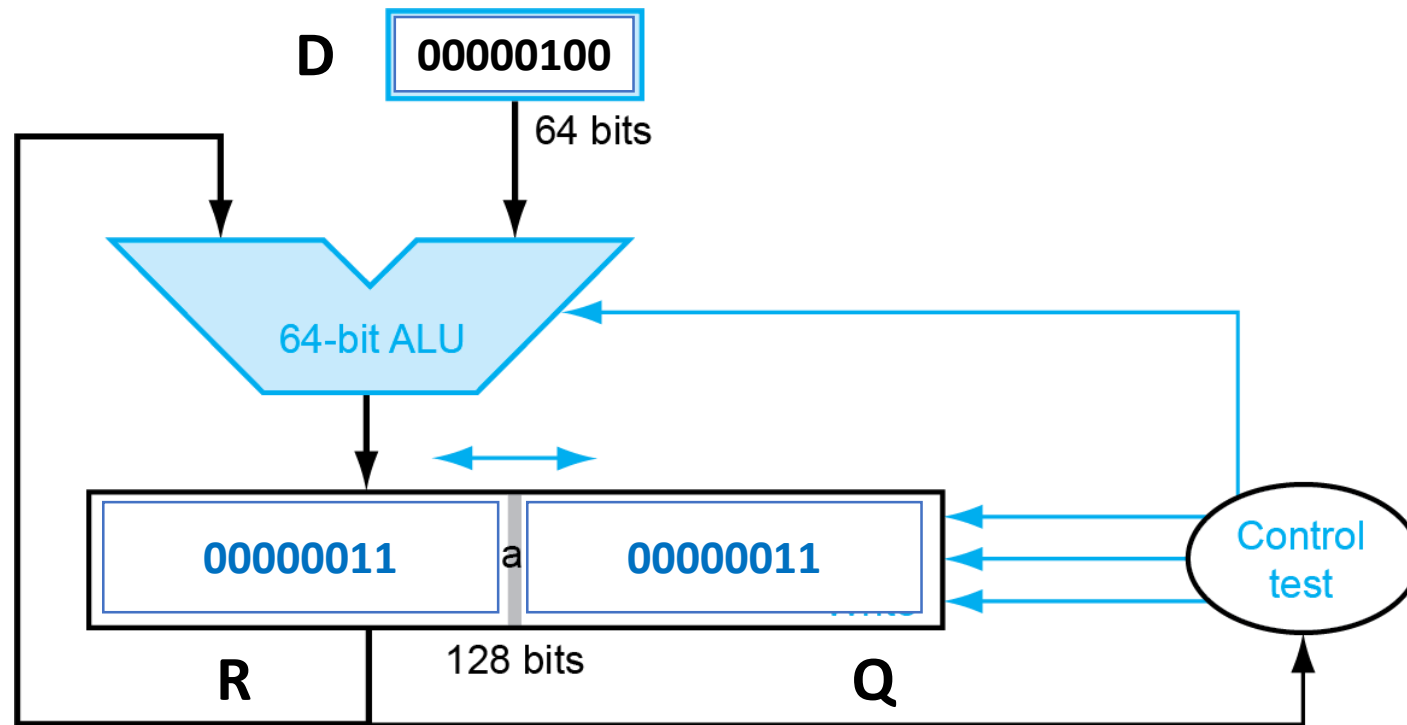
Division Example (4)

- $R - D \geq 0$: $R = R - D$, shift left R and Q by 1 bit, $Q_0 = 1$



Division Example (5)

- $R - D \geq 0$: $R = R - D$, shift left Q by 1 bit, $Q_0 = 1$



More on Division

■ Signed division

- Divide using absolute values
- Adjust sign of quotient and remainder as required

– (e.g.) $7 / 2$: $Q = 3, R = 1$
 $(-7) / 2$: $Q = -3, R = -1$
 $7 / (-2)$: $Q = -3, R = 1$
 $(-7) / (-2)$: $Q = 3, R = -1$

■ Faster division

- Can't use parallel hardware as in multiplier – Subtraction is conditional on sign of remainder
- Faster dividers (e.g. SRT division) generate multiple quotient bits per step, but still require multiple steps


$$-\left(\frac{x}{y}\right) == \frac{(-x)}{y}$$

$$\textit{Dividend} = \textit{Quotient} \times \textit{Divisor} + \textit{Remainder}$$

Shift for Multiplication/Division

■ Power-of-2 multiplication with left shift

- $u \ll k$ gives $u * 2^k$
 - e.g. $u \ll 3 == u * 8$, $(u \ll 5) - (u \ll 3) == u * 24$
- Both signed and unsigned
- Faster than multiply



■ Power-of-2 division with right shift

- $u \gg k$ gives $\lfloor u / 2^k \rfloor$
- Only for unsigned integers
 - e.g. $10_{10} / 4_{10} = 2_{10}$, $00001010_2 \gg 2 = 00000010_2 = 2_{10}$
 - $-10_{10} / 4_{10} = -2_{10}$, $11110110_2 \gg 2 = 11111101_2 = -3_{10}$
- For negative numbers, should be computed as $(u + (1 \ll k) - 1) \gg k$
 - e.g. $11110110_2 + 00000011_2 = 11111001_2$, $11111001_2 \gg 2 = 11111101_2 = -2_{10}$

Summary

- Two's complement for representing negative numbers
- Sign extension preserves the value
- Same hardware can be used for both signed and unsigned addition. This is also true for multiplication. But their overflow conditions are different.
- Remember?

```
int x = 50000;  
printf ("%s\n", (x*x >= 0)? "Yes" : "No");
```

- Using shift right for power-of-2 division requires correction for negative numbers