

Jin-Soo Kim
(jinsoo.kim@snu.ac.kr)

Systems Software &
Architecture Lab.

Seoul National University

Fall 2019

4190.308:

Computer Architecture



Course Information

- Schedule
 - 9:30 – 10:45 (Tuesday & Thursday) – **Also several supp. classes on Friday nights**
 - Lecture room: Engineering Bldg. #301-203
 - 3 credits
 - Official language: English
- TA: Jae-Hoon Shim (mattjs@snu)
- SNU eTL system for exam/project scores
- <http://csl.snu.ac.kr> for announcements and lecture slides
- <http://sys.snu.ac.kr> for project submissions and automatic grading

About Me

- Jin-Soo Kim (김진수)
 - Professor @ CSE Dept.
 - Systems Software & Architecture Laboratory
 - Operating systems, storage systems, parallel and distributed computing, embedded systems, ...
- E-mail: jinsoo.kim@snu.ac.kr
- Tel: 02-880-7302
- Office: Engineering Bldg. #301-520 (office hours: Tuesday & Thursday)
But, space remodeling is in progress until the end of September!
- The best way to contact me is by email

Myths About This Course

- It's an introductory course
 - Introduction to Computers?
 - About 20% of students have dropped every semester
- It's all about hardware
 - It's about how to separate work between software and hardware, and about how to design the interface between them
- It's not relevant for software engineers
 - Writing good software requires understanding details of underlying implementation
- Who needs to know assembly language these days?
 - Well, you'll see

Prerequisites


- Prerequisites
 - Programming Practice (4190.103A) – C programming
 - **Logic Design (MI522.000700) – Must!**
 - Data Structure (MI522.000900) – Recommended
- You should be familiar with the followings:
 - Shells and basic Linux commands
 - **C** (and **Python!**) programming skills
 - Basic knowledge on digital circuits and systems
- Accessible Linux (Ubuntu 18.04.3 LTS or similar) or MacOS machine

Policies for Non-major Students

- Your course registration form (“초안지”) will be accepted only if ...
 - You have an experience on C programming and debugging on Linux (gcc/gdb) and
 - You are familiar with Python and
 - You have already taken the “Logic Design” course
- Other introductory CSE courses for non-major students:
 - MI522.000600 Computer Programming
 - MI522.000700 Logic Design
 - MI522.000900 Data structures
 - 4190.101 Discrete Mathematics
 - 4190.103 Programming Practice


New Textbook

- Computer Organization and Design:
The Hardware/Software Interface
(**RISC-V Edition**)
 - David A. Patterson and John L. Hennessy
(Turing Award Recipients in 2017)
 - First Edition
 - Morgan Kaufmann, 2017
 - <http://booksite.elsevier.com/9780128122754/>




Previous Textbook

- Computer Systems:
A Programmer's Perspective
 - Randal E. Bryant and David R. O'Hallaron
 - **Third Edition**
 - Pearson Education Limited, 2016
 - Based on x86-64
 - <http://csapp.cs.cmu.edu>



Reference

- Computer Architecture:
A Quantitative Approach
 - John L. Hennessy and David A. Patterson
 - Sixth Edition
 - Morgan Kaufmann, 2017
 - <http://booksite.elsevier.com/9780128119051>



Topics

- Introduction to Computer Architecture
- Integers and Floating Points
- RISC-V Instruction Set Architecture
- Sequential Architecture
- Pipelined Architecture
- Cache
- Virtual memory
- I/O
- Parallel Computer Architecture

Projects (subject to change)

- C programming
- RISC-V assembly programming
- Designing pipelined processor
- Optimizing RISC-V assembly programs for pipelined processor
- Cache simulation

Grading Policy (subject to change)

- Exams: 60%
 - Midterm: 25%
 - Final: 35%
- Projects: 40%
- University policy requires students to attend at least 2/3 of the scheduled classes. Otherwise, you'll fail this course.
- We are using the electronic attendance system.
- Also, if you miss one of the exams, you'll fail this course.

Cheating Policy


- What is cheating?
 - Copying another student's solution (or one from the Internet) and submitting it as your own
 - Allowing another student to copy your solution
- What is NOT cheating?
 - Helping others use systems or tools
 - Helping others with high-level design issues
 - Helping others debug their code
- Penalty for cheating
 - Severe penalty on the grade (F) and report to dept. chair
 - Ask helps to your TA or instructor if you experience any difficulty!

Introduction to Computer Architecture

Classes of Computers


- **Personal computers**
 - General purpose, variety of software
 - Subject to cost/performance tradeoff
- **Supercomputers**
 - High-end scientific and engineering calculations
 - Highest capability but represent a small fraction of the overall computer market
- **Server computers**
 - Network based
 - High capacity, performance, reliability
 - Range from small servers to large data centers
- **Embedded computers**
 - Hidden as components of systems
 - Stringent power/performance/cost constraints

Opening the Box (iMac Pro)



Components of a Computer

- CPU: Control + Datapath
- Memory
- I/Os
 - User-interface devices:
Display, keyboard, mouse, sound, ...
 - Storage devices:
HDD, SSD, CD/DVD, ...
 - Network adapters:
Ethernet, 3G/4G/5G, WiFi, Bluetooth, ...
- Same components for all kinds of computer



What Happened:

1997

2017

104 cabinets
(76 computes,
8 switches,
20 disks)

9298 cores

150m²



ASCI Red at Sandia

1.3 TF/s, 850 KW



Cavium ThunderX2

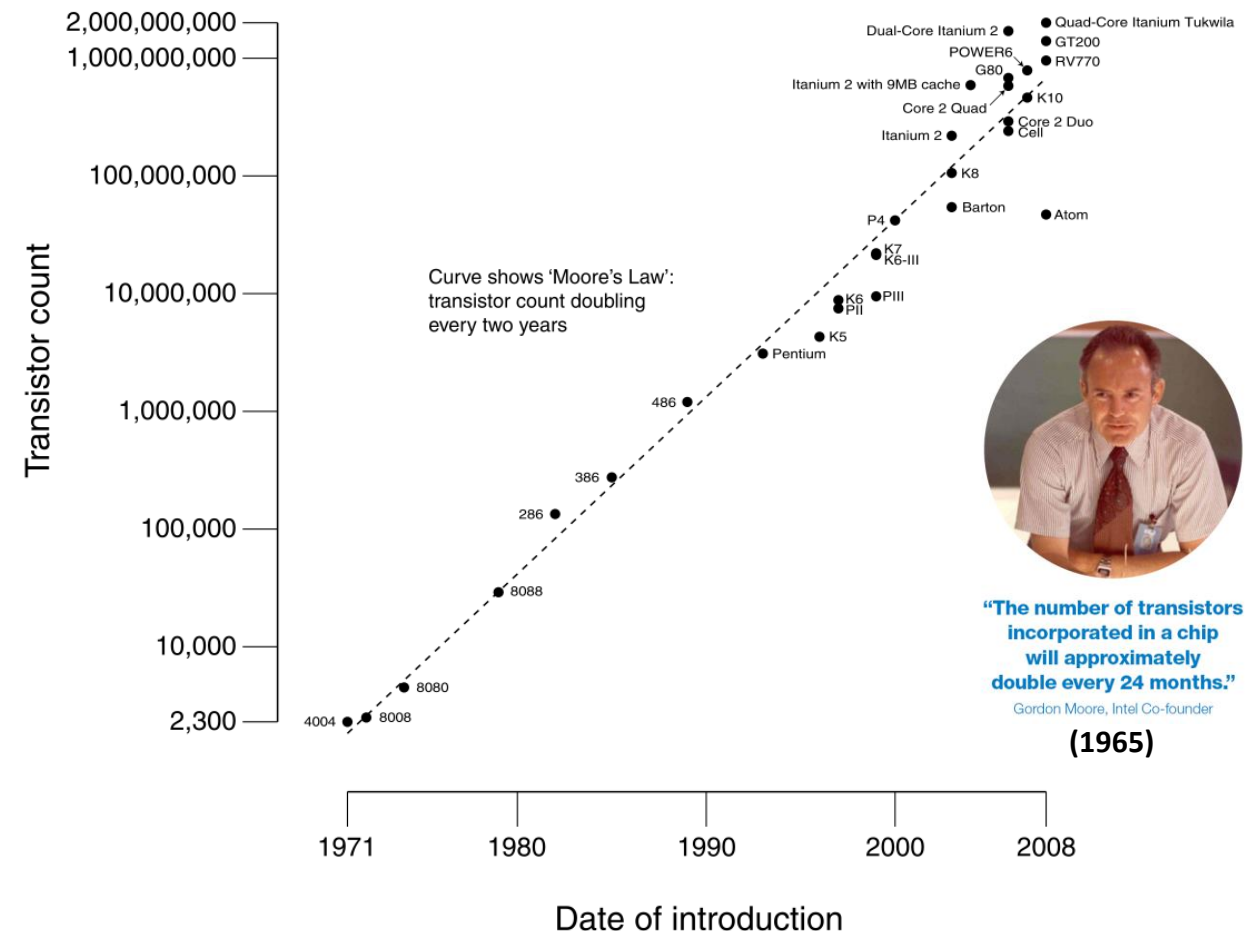
~ 1.1 TF/s, ~ 0.2 KW

3.5 orders of
magnitude

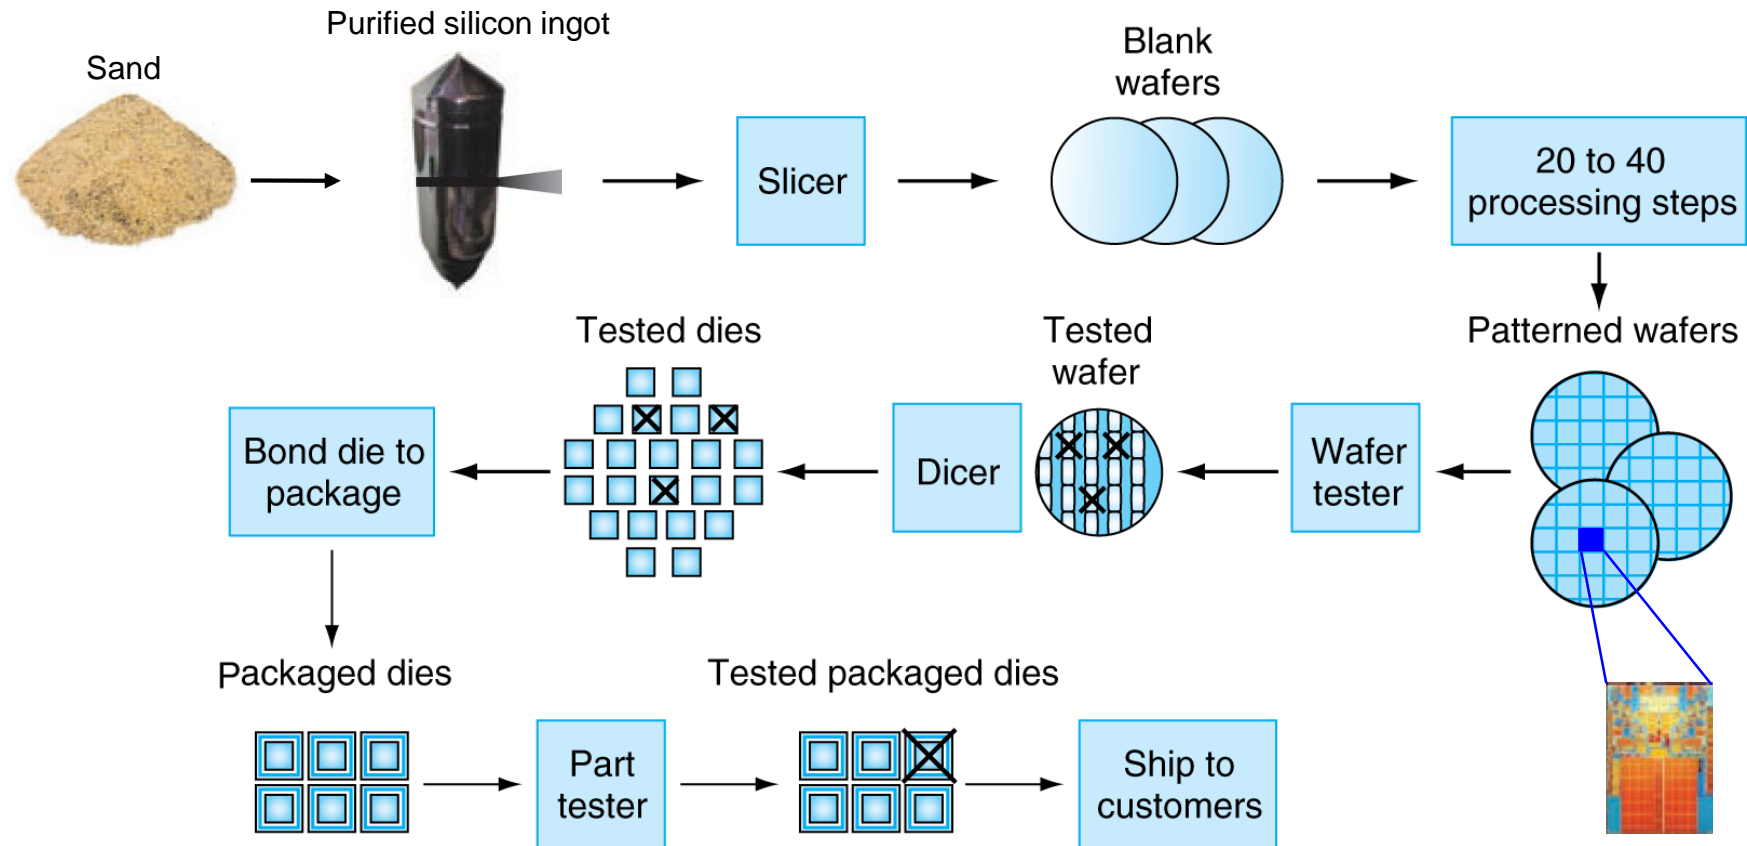
The Computer Revolution

- Progress in computer technology
 - Underpinned by **Moore's Law**
- Makes novel applications feasible
 - World Wide Web (WWW)
 - Smartphones
 - Search engines
 - Human genome project
 - Self-driving cars
 - Artificial intelligence
 - VR/AR
- Computers are pervasive

CPU Transistor Counts 1971-2008 & Moore's Law



From Sand to Circuits




Yield: proportion of working dies per wafer


Processing ICs




1. Start with a partially processed die on a silicon wafer.




2. Deposit oxide layer.




3. Coat with photoresist.




4. Position mask and flash ultraviolet light.



5. Rinse with solvent.






6. Etch with acid.



7. Remove remaining photoresist.

Realizing a Logic Gate

- NAND logic built with CMOS technology




Intel Core i7 Wafer

- 12-inch (300mm) wafer, 280 chips, 32nm technology
- Each chip is 20.7 x 10.5 mm

$$\text{Cost per die} = \frac{\text{Cost per wafer}}{\text{Dies per wafer} \times \text{Yield}}$$

$$\text{Dies per wafer} \approx \text{Wafer area} / \text{Die area}$$

$$\text{Yield} = \frac{1}{\left(1 + \left(\text{Defects per area} \times \frac{\text{Die area}}{2}\right)\right)^2}$$



- Wafer cost and area are fixed
- Defect rate determined by manufacturing process
- Die area determined by architecture and circuit design

Technology Trends


■ CPU

- Logic capacity: $\sim 30\%$ / year
- Clock rate: $\sim 20\%$ / year

■ Memory

- DRAM capacity: $\sim 60\%$ / year
(4x every 3 years)
- DRAM speed: $\sim 10\%$ / year


Year	Technology	Relative performance/cost
1951	Vacuum tube	1
1965	Transistor	35
1975	Integrated circuit (IC)	900
1995	Very large scale IC (VLSI)	2,400,000
2013	Ultra large scale IC	250,000,000,000



Architecture: Exploiting Parallelism

- **Instruction level parallelism (ILP)**
 - Pipelining
 - Superscalar
 - Out-of-order execution
 - Branch prediction
 - VLIW (Very Long Instruction Word)
- **Data level parallelism (DLP)**
 - SIMD / Vector instructions
- **Task level parallelism (TLP)**
 - Simultaneous multithreading (Hyperthreading)
 - Multicore


Uniprocessor Performance



Power Trends

- In CMOS IC technology,

$$Power = Capacitive Load \times Voltage^2 \times Frequency$$



Reducing Power


- Suppose a new CPU has
 - 85% of capacitive load of old CPU
 - 15% voltage and 15% frequency reduction

$$\frac{P_{new}}{P_{old}} = \frac{C_{old} \times 0.85 \times (V_{old} \times 0.85)^2 \times F_{old} \times 0.85}{C_{old} \times V_{old}^2 \times F_{old}} = 0.85^4 = 0.52$$

- The power wall
 - We can't reduce voltage further
 - We can't remove more heat
- How else can we improve performance?

The Shift to Multicores

- ILP wall
 - Control dependency
 - Data dependency
- Memory wall
 - Memory latency improved by 10% / year
 - Cache shows diminishing returns
- Power wall




Amdahl's Law


- The theoretical upper limit of speed up is limited by the serial portion of the code

$$\text{Speedup} = \frac{1}{(1 - P) + \frac{P}{n}}$$

- $S = (1 - P)$: the time spent executing the serial portion
 - n : the number of processor cores
-
- Corollary: make the common case fast




A New Golden Age?



**End of
Moore's law**

**A New Golden Age
for Computer
Architecture
(CACM, Feb. 2019)**

Intel Core i9-9900K (Coffee Lake, 2018)

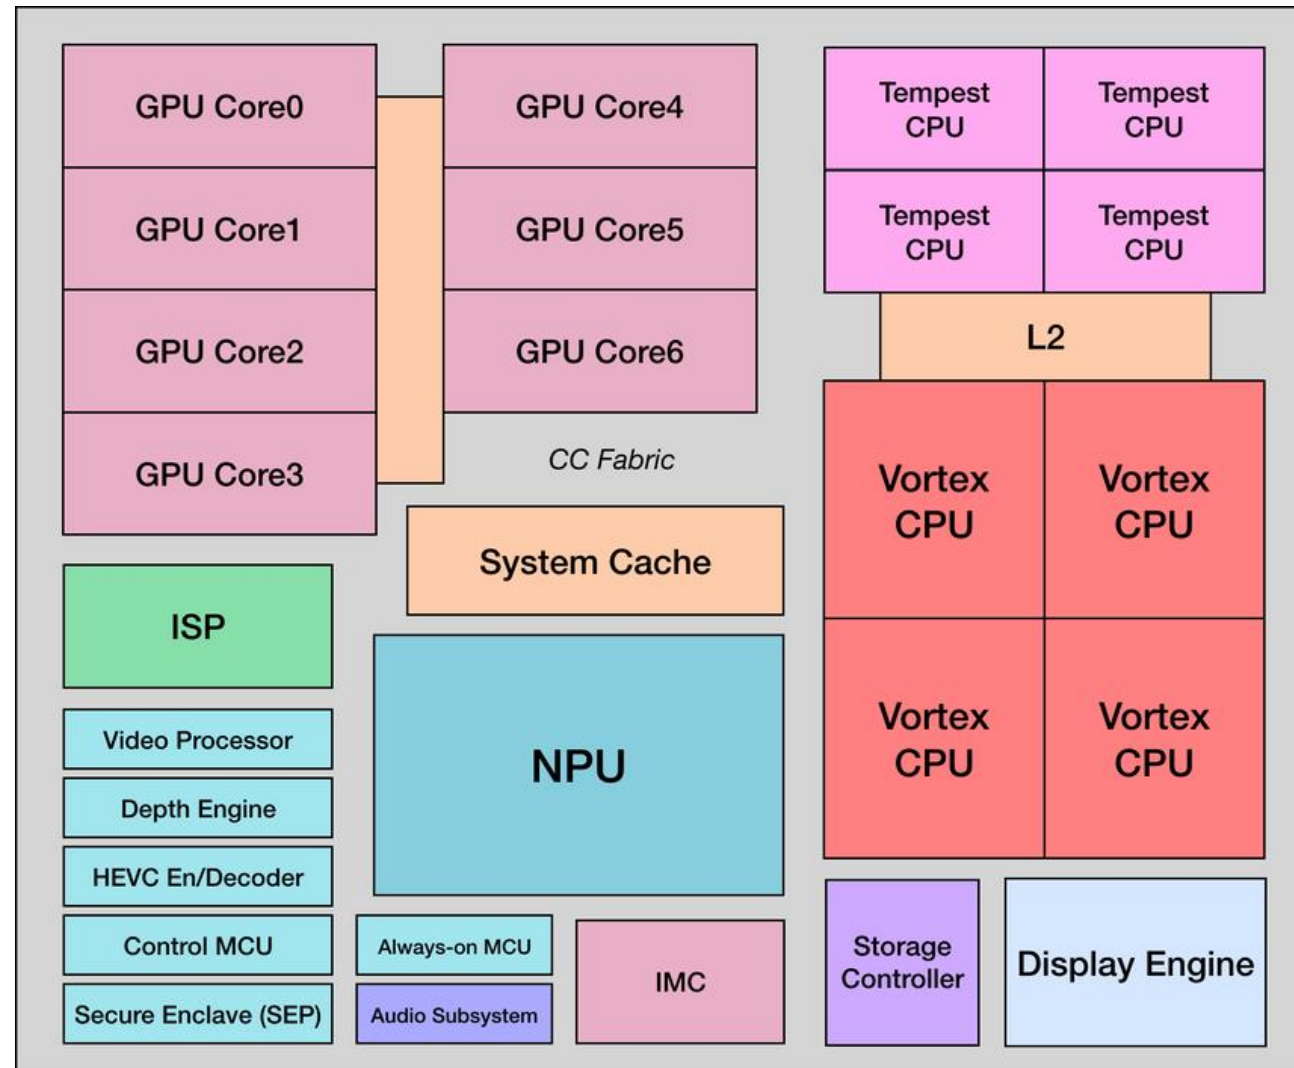


Process:
14nm

Transistors:
~ 3B

Die size:
~ 177 mm²

Apple A12X Bionic (2018)



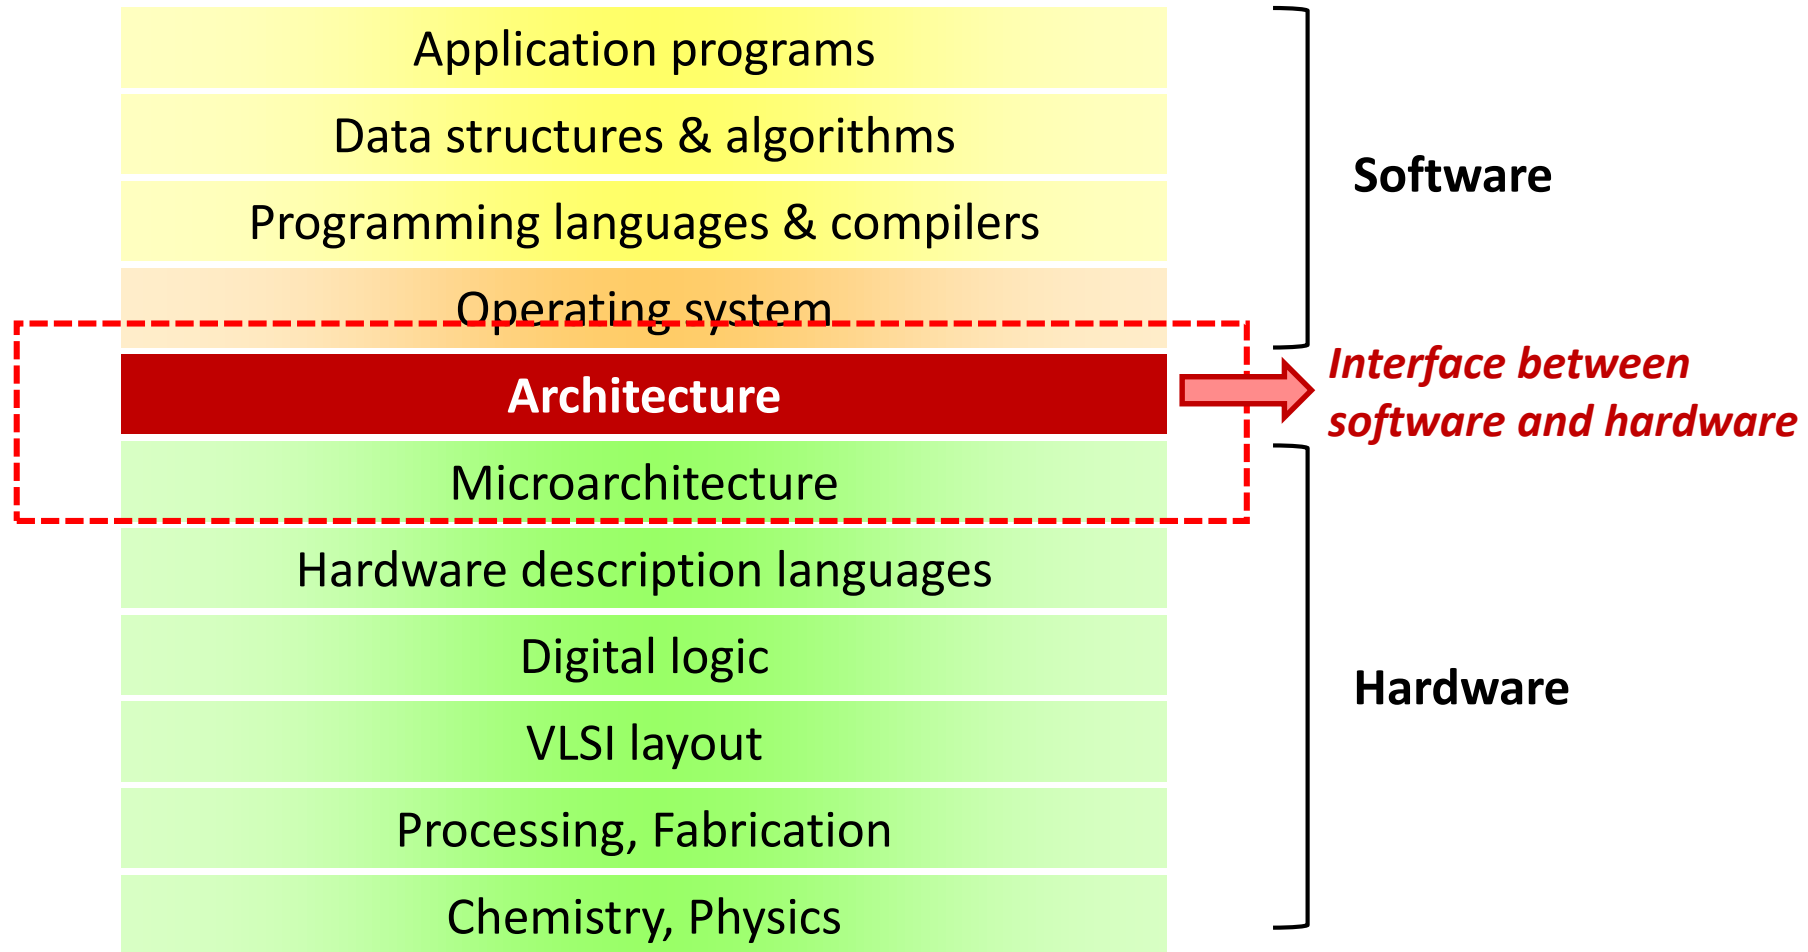
Process:
7nm

Transistors:
~ 10B

Die size:
~ 122 mm²


Taming Complexity

- Levels of abstractions




Below Your Program

- **Application software**
 - Written in high-level language
- **System software**
 - Compiler: translates HLL code to machine code
 - Operating system: service code
 - Handling input/output
 - Managing memory and storage
 - Scheduling tasks & sharing resources
- **Hardware**
 - Processor, memory, I/O controllers



Instruction Set Architecture (ISA)

- The hardware/software interface
 - Hardware abstraction visible to software (OS, compilers, ...)
 - Instructions and their encodings, registers, data types, addressing modes, etc.
 - Written documents about how the CPU behaves
 - e.g. All 64-bit Intel CPUs follow the same x86-64 (or Intel 64) ISA



Levels of Program Code

- **High-level language**
 - Level of abstraction closer to problem domain
 - Provides for productivity and portability
- **Assembly language**
 - Textual representation of instructions
 - For humans
- **Machine language**
 - Hardware representation
 - Binary digits (bits)
 - Encoded instructions and data

High-level
language
program
(in C)

```
swap(int v[], int k)
{int temp;
  temp = v[k];
  v[k] = v[k+1];
  v[k+1] = temp;
}
```

Compiler

Assembly
language
program
(for RISC-V)

```
swap:
    slli x6, x11, 3
    add  x6, x10, x6
    ld   x5, 0(x6)
    ld   x7, 8(x6)
    sd   x7, 0(x6)
    sd   x5, 8(x6)
    jalr x0, 0(x1)
```

Assembler

Binary machine
language
program
(for RISC-V)

```
00000000001101011001001100010011
00000000011001010000001100110011
00000000000000110011001010000011
00000000100000110011001110000011
00000000011100110011000000100011
00000000010100110011010000100011
0000000000000001000000001100111
```

Abstraction is Good, But ...

- Abstraction helps us deal with complexity
 - Hide lower-level details
- These abstractions have limits
 - Especially in the presence of bugs
 - Need to understand details of underlying implementations
- This is why you should take this course seriously even if you don't want to be a computer architect!

Example #1: Int's \neq Integers, Float's \neq Reals

■ Is $x^2 \geq 0$?

- Float's: ??
- Int's: ??

```
int x = 50000;  
printf ("%s\n", (x*x >= 0)? "Yes" : "No");
```

■ Is $(x + y) + z == x + (y + z)$?

- Unsigned & Signed Int's: ??
- Float's: ??

```
float x = 1e20, y = -1e20, z = 3.14;  
printf ("%s\n", (x+y)+z==x+(y+z)? "Yes" : "No");
```

Example #2: More Than Just GHz

CPU	Clock Speed	SPECint2000	SPECfp2000
Athlon 64 FX-55	2.6GHz	1854	1782
Pentium 4 Extreme Edition	3.46GHz	1772	1724
Pentium 4 Prescott	3.8GHz	1671	1842
Opteron 150	2.4GHz	1655	1644
Itanium 2 9MB	1.6GHz	1590	2712
Pentium M 755	2.0GHz	1541	1088
POWER5	1.9GHz	1452	2702
SPARC64 V	1.89GHz	1345	1803
Athlon 64 3200+	2.2GHz	1080	1250
Alpha 21264C	1.25GHz	928	1019

Example #3: Constant Factors Matter

- There's more to performance than asymptotic complexity
- Array copy example

```
void copyij (int src[2048][2048],  
             int dst[2048][2048])  
{  
    int i, j;  
    for (i = 0; i < 2048; i++)  
        for (j = 0; j < 2048; j++)  
            dst[i][j] = src[i][j];  
}
```

4.3 ms

```
void copyji (int src[2048][2048],  
             int dst[2048][2048])  
{  
    int i, j;  
    for (j = 0; j < 2048; j++)  
        for (i = 0; i < 2048; i++)  
            dst[i][j] = src[i][j];  
}
```

81.8 ms

copyji() is 20x slower on 2.0GHz Intel Core i7 Haswell. Why?

Example #4: Memory Matters

- Memory referencing bug example

```
/* Echo Line */
void echo()
{
    // Way too small!
    char buf[4];
    gets(buf);
    puts(buf);
}

int main()
{
    printf("Type: ");
    echo();
    return 0;
}
```

```
$ ./bufdemo
```

```
Type:012
```

```
012
```

```
$ ./bufdemo
```

```
Type: 01234567890123456789012
```

```
01234567890123456789012
```

```
$ ./bufdemo
```

```
Type: 012345678901234567890123
```

```
Segmentation fault (core dumped)
```

What You Will Learn

- How data are represented?
- How programs are translated into the machine language
 - And how the hardware executes them
- The hardware/software interface – Instruction Set Architecture (ISA)
- What determines program performance
- How hardware designers / software developers improve performance
- What is parallel processing

Eight Great Ideas in Computer Architecture

- Design for **Moore's Law**
- Use **abstraction** to simplify design
- Make the **common case fast**
- Performance via **parallelism**
- Performance via **pipelining**
- Performance via **prediction**
- **Hierarchy** of memories
- **Dependability** via redundancy



Why Take This Course?

- To graduate!
- To design the next great instruction set? Well...
 - ISA has largely converged, especially in desktop / server / laptop / mobile space
 - Dictated by powerful market forces (Intel/ARM)
- To get a job in Intel, NVIDIA, ARM, Apple, Qualcomm, Google, ...
 - Tremendous organizational innovations relative to established ISA abstractions
- Design, analysis, and implementation concepts that you'll learn are vital to all aspects of computer science and engineering
- This course will equip you with an intellectual toolbox for dealing with a host of systems design challenges
- And finally, just for fun!

Summary

- Modern Computer Architecture is about managing and optimizing across several levels of abstraction w.r.t. dramatically changing technology and application load
- This course focuses on
 - RISC-V Instruction Set Architecture (ISA) – A new open interface
 - An implementation based on Pipelining (Microarchitecture) – how to make it faster?
 - Memory hierarchy – How to make trade-offs between performance and cost?
- Understanding Computer Architecture is vital to other “systems” courses:
 - System programming, Operating systems, Compilers, Embedded systems, Computer networks, Multicore computing, Distributed systems, Mobile computing, Security, Machine learning, Quantum computing, etc.